

Unit 2 Lecture 3: Cross-validation

September 22, 2022

In this R demo, we will implement cross-validation to select the degrees of freedom of a natural spline fit, using the running example from the previous class. We'll need the following R packages:

```
library(tidyverse)
library(splines)    # for ns()
library(cowplot)    # for plot_grid()
library(stat471)    # for cross_validate_spline()
```

Training and validation

Let us create a training set:

```
set.seed(1)
f <- function(x) (sin(3 * x))
n <- 50
sigma <- 1
train_data <- tibble(
  x = seq(0, 2 * pi, length.out = n),
  y = f(x) + rnorm(n, sd = sigma)
)
train_data
```

```
## # A tibble: 50 x 2
##       x      y
##   <dbl> <dbl>
## 1 0      -0.626
## 2 0.128  0.559
## 3 0.256 -0.140
## 4 0.385  2.51
## 5 0.513  1.33
## 6 0.641  0.118
## 7 0.769  1.23
## 8 0.898  1.17
## 9 1.03   0.640
## 10 1.15  -0.620
## # ... with 40 more rows
## # i Use `print(n = ...)` to see more rows
```

Let's also suppose we have a large validation set on our hands:

```
N <- 50000
validation_data <- tibble(
  x = seq(0, 2 * pi, length.out = N),
  y = f(x) + rnorm(n, sd = sigma)
)
```

Now let's fit splines with $df = 1, 2, \dots, 15$ to the training data, and evaluate their test error using the test set:


```

# compute the validation error
max_df <- 15
validation_error <- numeric(max_df)
df <- 1
for (df in 1:max_df) {
  formula <- sprintf("y ~ ns(x, df = %d)", df)
  spline_fit <- lm(formula = formula, data = train_data)
  y_hat_validation <- predict(spline_fit, newdata = validation_data)
  validation_error[df] <- validation_data %>%
    cbind(y_hat_validation) %>%
    summarise(mean((y_hat_validation - y)^2)) %>%
    pull()
}
validation_error

## [1] 1.3863485 1.3893519 1.3683364 1.3577412 1.2534633 1.3230281 0.9946587
## [8] 1.0148754 0.9779877 1.0783341 1.0807794 1.0606353 1.0973088 1.1098777
## [15] 1.1127368

# plot the validation error
p_val <- tibble(df = 1:max_df, validation_error) %>%
  ggplot(aes(x = df, y = validation_error)) +
  geom_point() +
  geom_line() +
  labs(
    x = "Degrees of freedom",
    y = "Validation error"
  )
plot(p_val)

```



The issue is that we usually do not have a giant validation set for model selection purposes. We need to make do with our smallish training set for both model training and model selection. This is where cross-validation comes in handy!

Cross-validation for $df = 5$

The idea is to split our training samples into *folds* and then have the folds take turns being the validation set. Let's take a look.

```
K <- 10
folds <- sample(rep(1:K, n / K))
train_data <- train_data %>%
  mutate(fold = folds)
train_data

## # A tibble: 50 x 3
##       x       y fold
##   <dbl> <dbl> <int>
## 1 0      -0.626   3
## 2 0.128  0.559   6
## 3 0.256 -0.140   4
## 4 0.385  2.51    2
## 5 0.513  1.33    9
## 6 0.641  0.118   7
## 7 0.769  1.23    8
## 8 0.898  1.17    1
## 9 1.03   0.640   1
## 10 1.15  -0.620   5
## # ... with 40 more rows
## # i Use `print(n = ...)` to see more rows
```

Question: How would we select the data in fold number 1? How would we select all the data except fold number 1?

Let's first use cross-validation to estimate the test error for a spline fit with 5 degrees of freedom.

```
# create a vector of out-of-fold predictions
out_of_fold_predictions <- numeric(n)

# iterate over folds
for (current_fold in 1:K) {
  # out-of-fold data will be used for training
  out_of_fold_data <- train_data %>% filter(fold != current_fold)
  # in-fold data will be used for validation
  in_fold_data <- train_data %>% filter(fold == current_fold)

  out_of_fold_data
  in_fold_data

  # train on out-of-fold data
  spline_fit <- lm(y ~ ns(x, df = 5), data = out_of_fold_data)

  # predict on in-fold data
  out_of_fold_predictions[folds == current_fold] <-
    predict(spline_fit, newdata = in_fold_data)
}

# add the out-of-fold predictions to the data frame
results <- train_data %>%
  mutate(yhat = out_of_fold_predictions)
```

```
results
```

```
## # A tibble: 50 x 4
##       x       y fold  yhat
##   <dbl> <dbl> <int> <dbl>
## 1 0      -0.626   3  1.85
## 2 0.128  0.559   6  0.726
## 3 0.256 -0.140   4  0.878
## 4 0.385  2.51    2  0.252
## 5 0.513  1.33    9  0.242
## 6 0.641  0.118   7  0.181
## 7 0.769  1.23    8 -0.0212
## 8 0.898  1.17    1 -0.374
## 9 1.03   0.640   1 -0.490
## 10 1.15  -0.620   5  0.0871
## # ... with 40 more rows
## # i Use `print(n = ...)` to see more rows
```

```
# compute the CV estimate and standard error
results %>%
  group_by(fold) %>%
  summarise(cv_fold = mean((yhat - y)^2)) %>% # CV estimates per fold
  summarise(
    cv_mean = mean(cv_fold),
    cv_se = sd(cv_fold) / sqrt(K)
  )
```

```
## # A tibble: 1 x 2
##   cv_mean cv_se
##   <dbl> <dbl>
## 1    1.70 0.339
```

What are two reasons this CV estimate may be different from the validation error estimated above?

Cross-validation for $df = 1, 2, \dots, 15$

Now let's repeat what we did above for many degrees of freedom, because after all, the point of cross-validation is to choose the degrees of freedom. We can do this via `cross_validate_spline()` in the `stat471` package.

```
cv_results <- cross_validate_spline(
  x = train_data$x,
  y = train_data$y,
  nfolds = 10,
  df_values = 1:15
)
```

The `cv_results` variable stores an *object*, with the following fields:

```
names(cv_results)
```

```
## [1] "cv_table" "cv_plot" "df.1se" "df.min"
```


Let's inspect these one at a time:

```
cv_results$cv_table
```

```
## # A tibble: 15 x 3
##       df cv_mean cv_se
```



```
##      <int>      <dbl> <dbl>
## 1         1      1.35 0.280
## 2         2      1.38 0.276
## 3         3      1.42 0.268
## 4         4      1.55 0.298
## 5         5      1.55 0.282
## 6         6      1.81 0.301
## 7         7      1.16 0.285
## 8         8      0.928 0.178
## 9         9      0.840 0.198
## 10        10      0.828 0.162
## 11        11      0.909 0.166
## 12        12      0.903 0.170
## 13        13      0.900 0.167
## 14        14      0.951 0.149
## 15        15      0.948 0.191
```

```
cv_results$cv_plot
```



Let's compare the CV plot to the validation error plot:

```
plot_grid(cv_results$cv_plot, p_val, nrow = 2)
```



Based on this plot, what degrees of freedom would we select using the one-standard error rule?

```
cv_results$df.1se
```

```
## [1] 8
```

```
cv_results$df.min
```

```
## [1] 10
```

Exercise

Use 10-fold cross-validation and the one-standard error rule to select the optimal number of degrees of freedom for regressing `income` on `age` in the `income_2007` data from Unit 3 Lecture 1. Make the CV plot, extract `df.1se` and produce a scatter plot of `income` versus `age` with this optimal spline fit superimposed.

```
income_2007 <- read_csv("income_data.csv") %>%
  filter(year == 2007) %>%
  select(-year)
```