

# Neural networks

STAT 4710

November 17, 2022

# Where we are

- ✓ **Unit 1:** R for data mining
- ✓ **Unit 2:** Prediction fundamentals
- ✓ **Unit 3:** Regression-based methods
- ✓ **Unit 4:** Tree-based methods
- Unit 5:** Deep learning

**Lecture 1:** Deep learning preliminaries

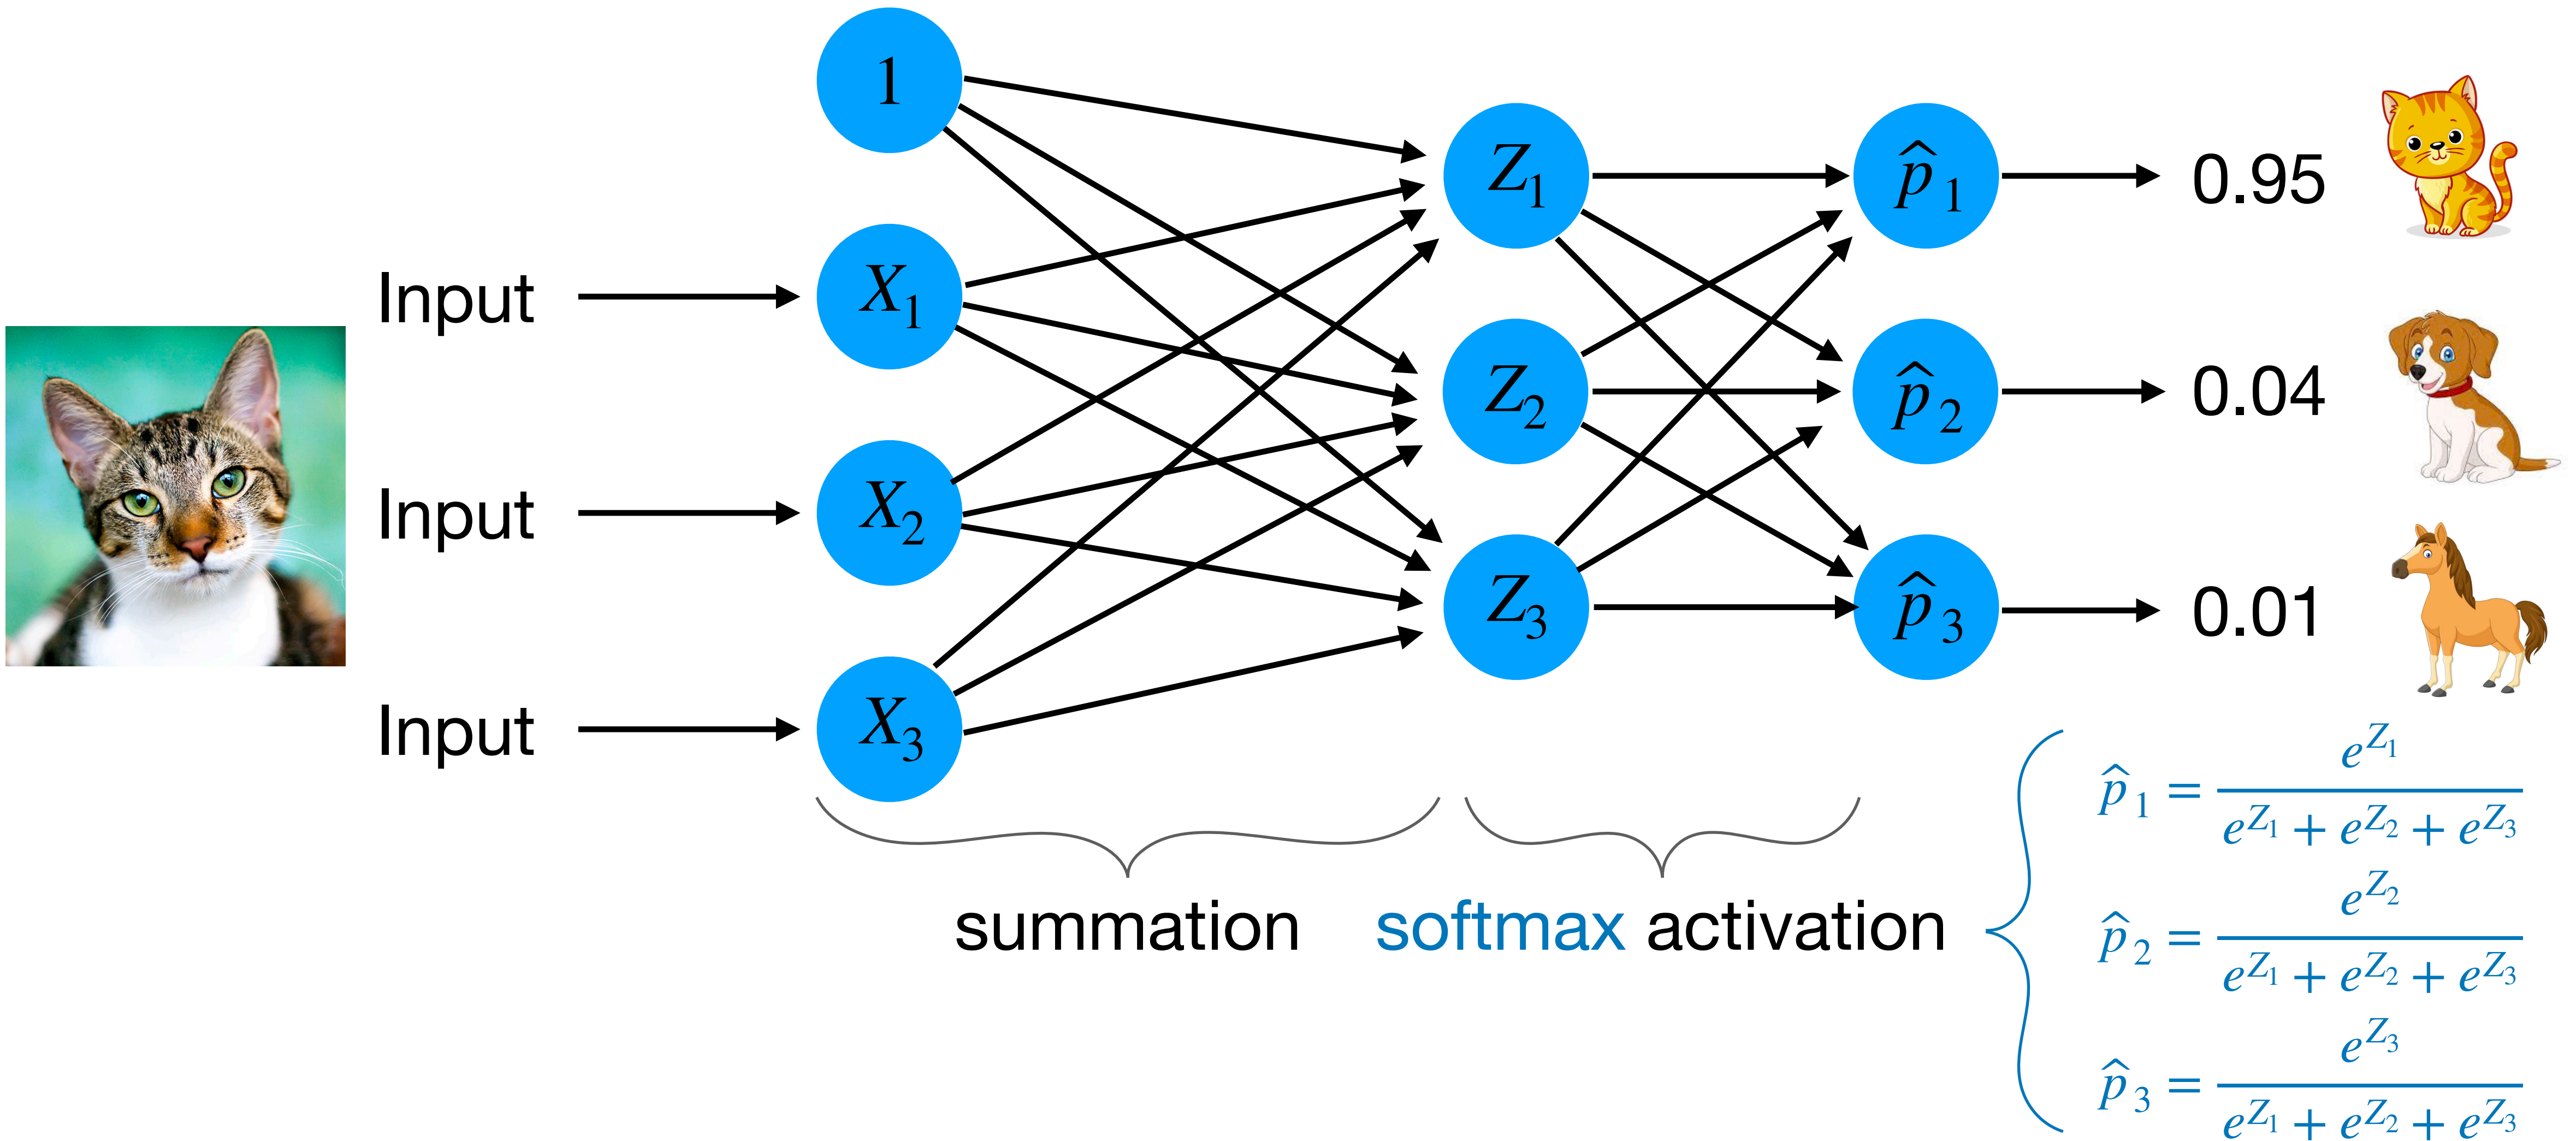
**Lecture 2:** Neural networks

**Lecture 3:** Deep learning for images

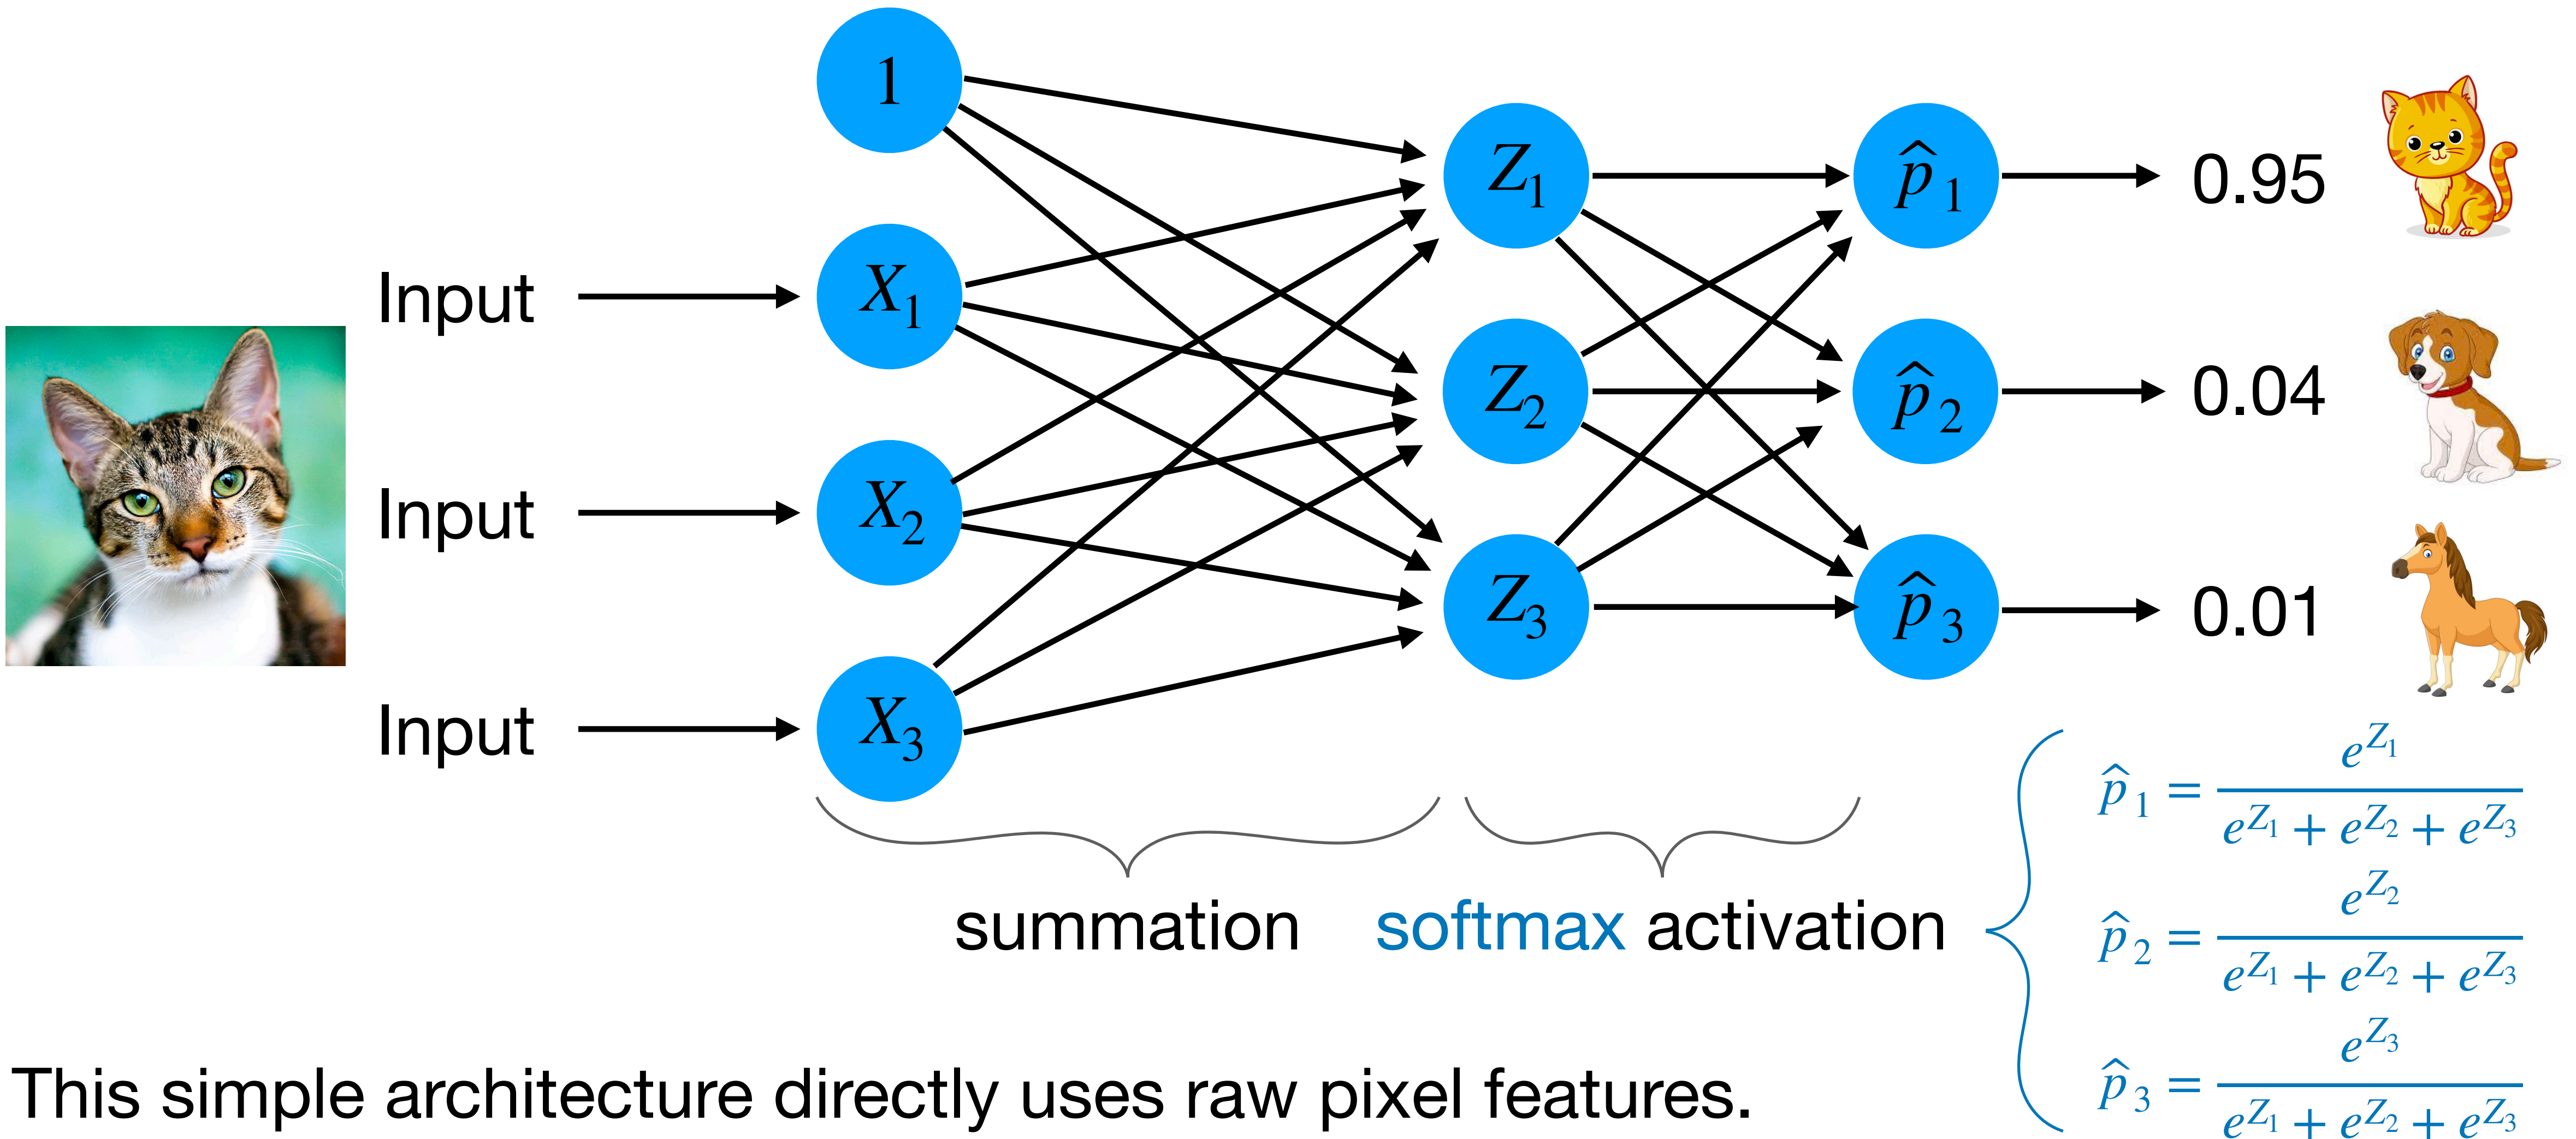
**Lecture 4:** Deep learning for text

**Lecture 5:** Unit review and quiz in class

# Recall: Multi-class logistic model



# Recall: Multi-class logistic model



This simple architecture directly uses raw pixel features.  
Intuitively, a good classifier would pick up on cat's eyes, ears, etc.

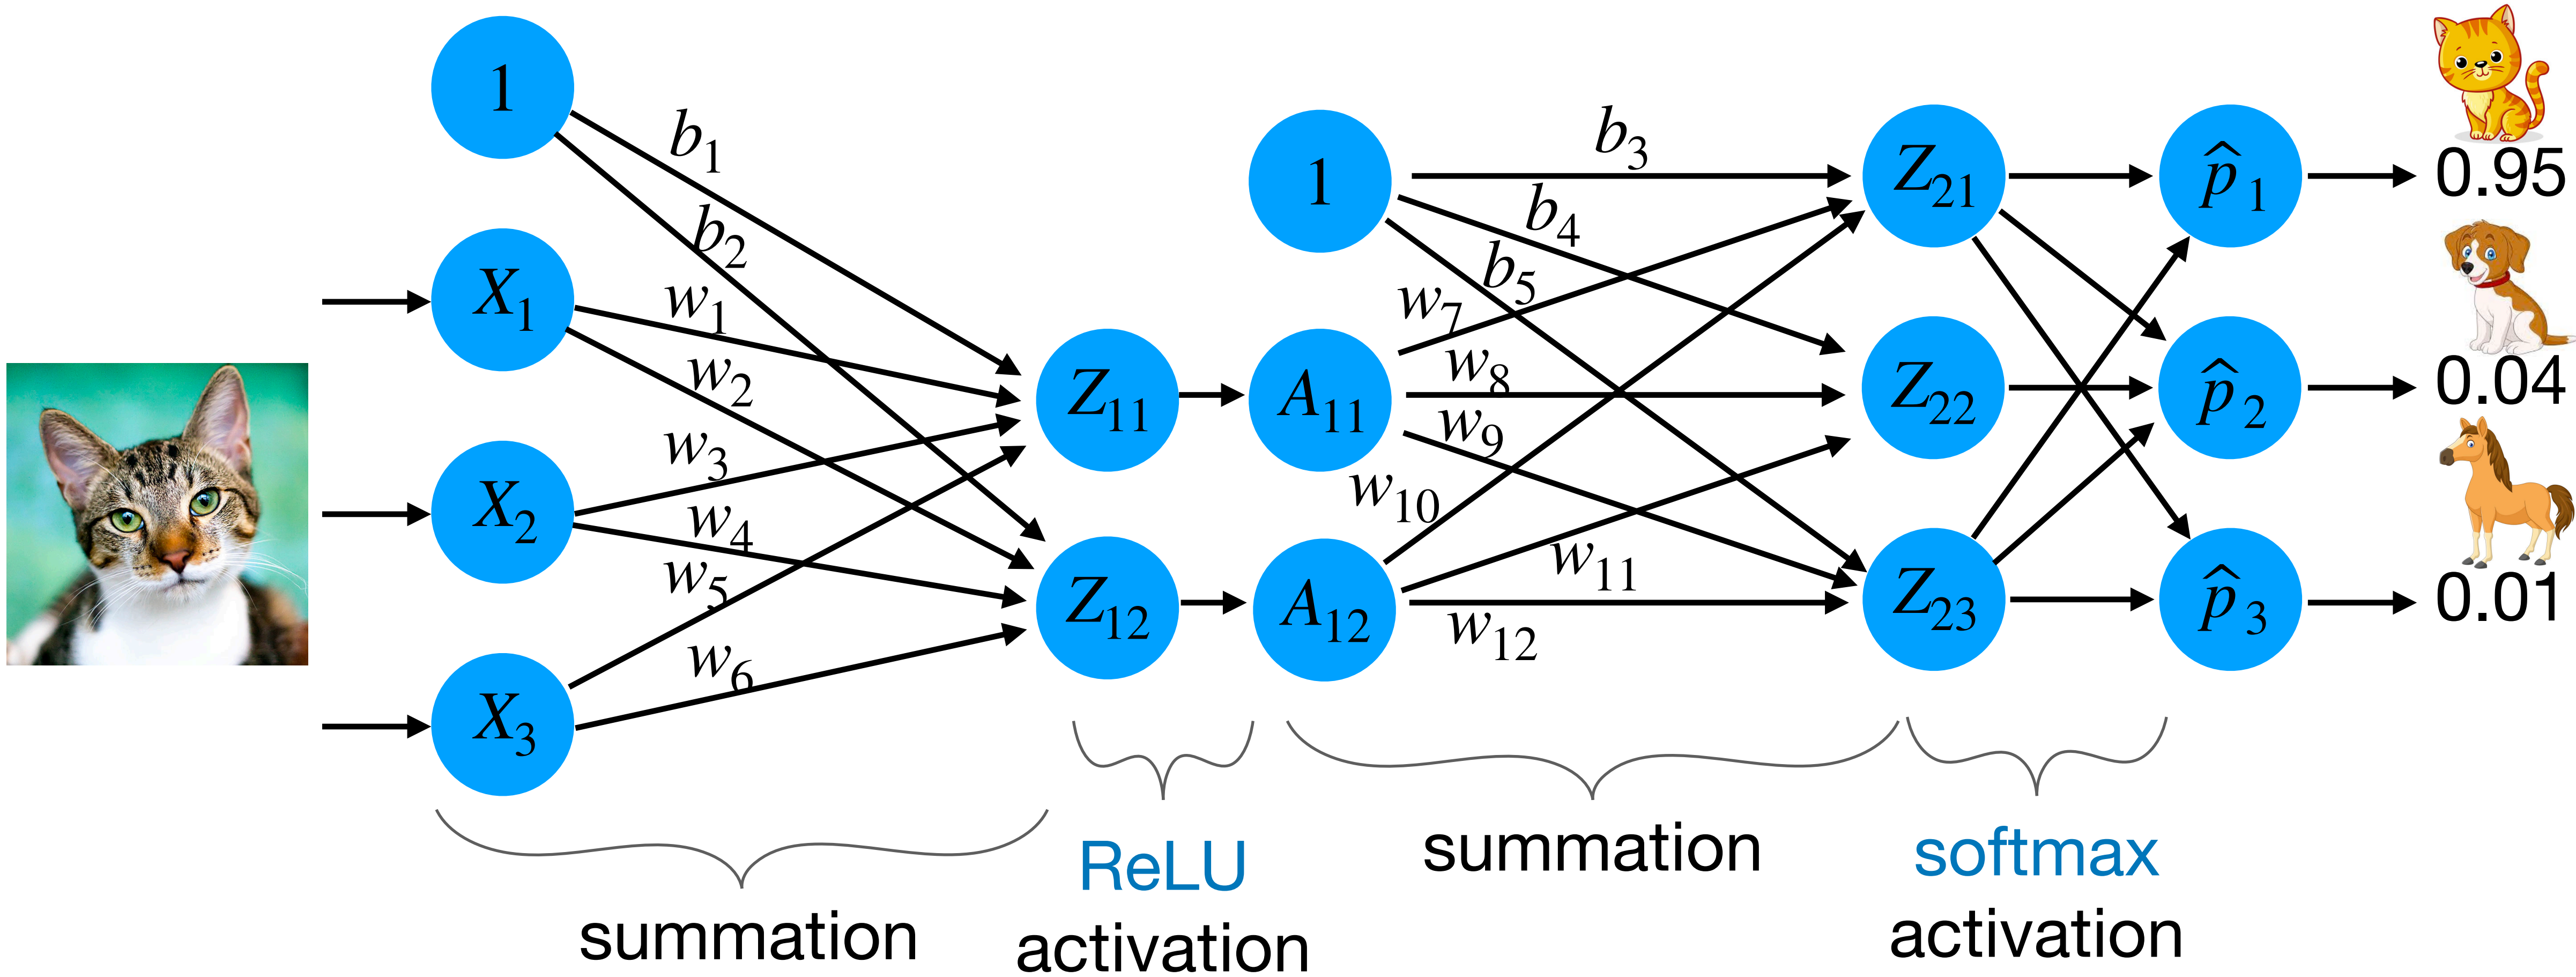


# One-hidden-layer neural network

Input layer

Hidden layer

Output layer

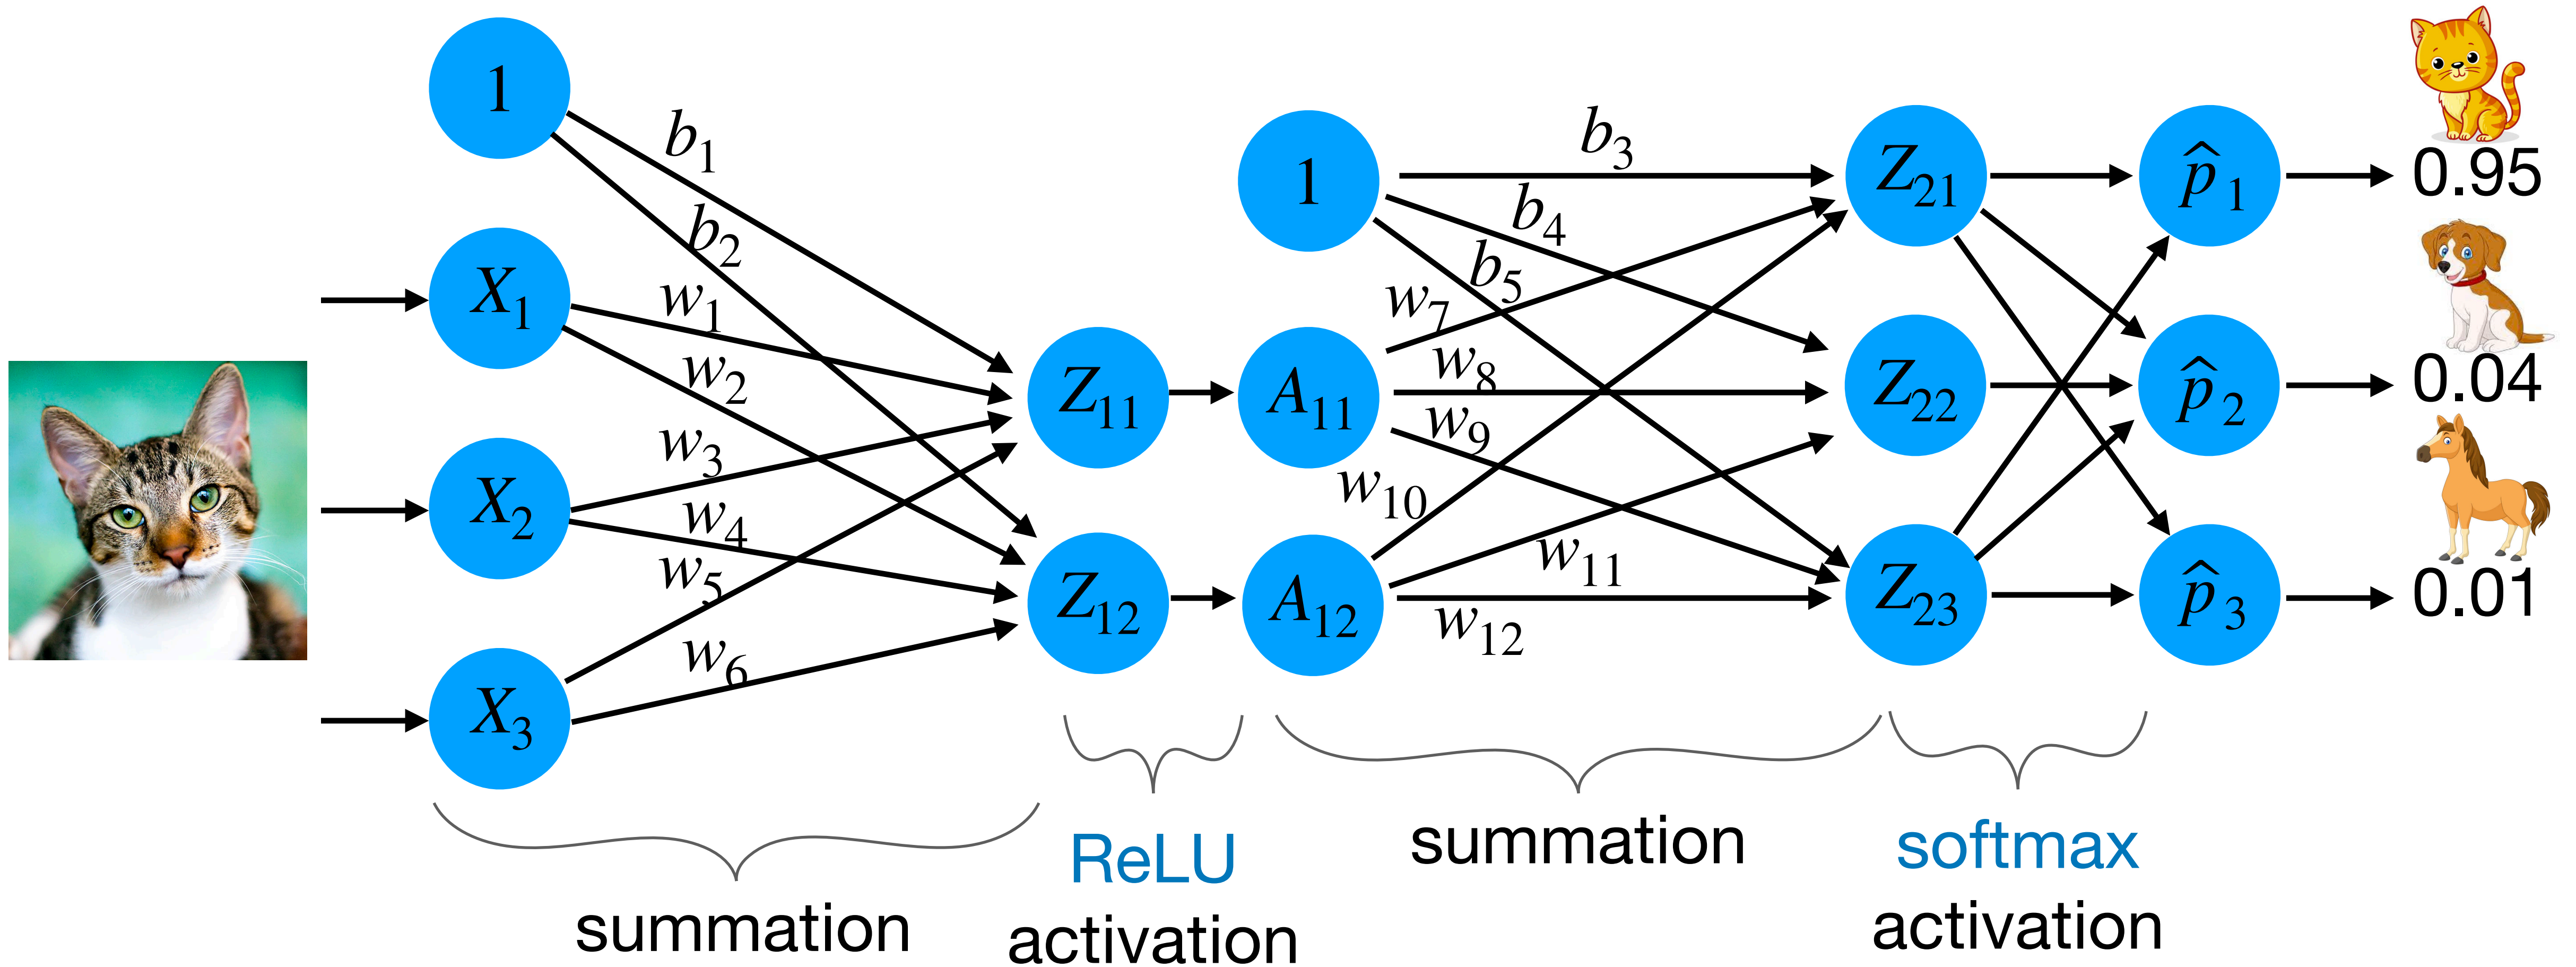


# One-hidden-layer neural network

Input layer

Hidden layer

Output layer



E.g.,  $Z_{12} = b_2 + w_2X_1 + w_4X_2 + w_6X_3$

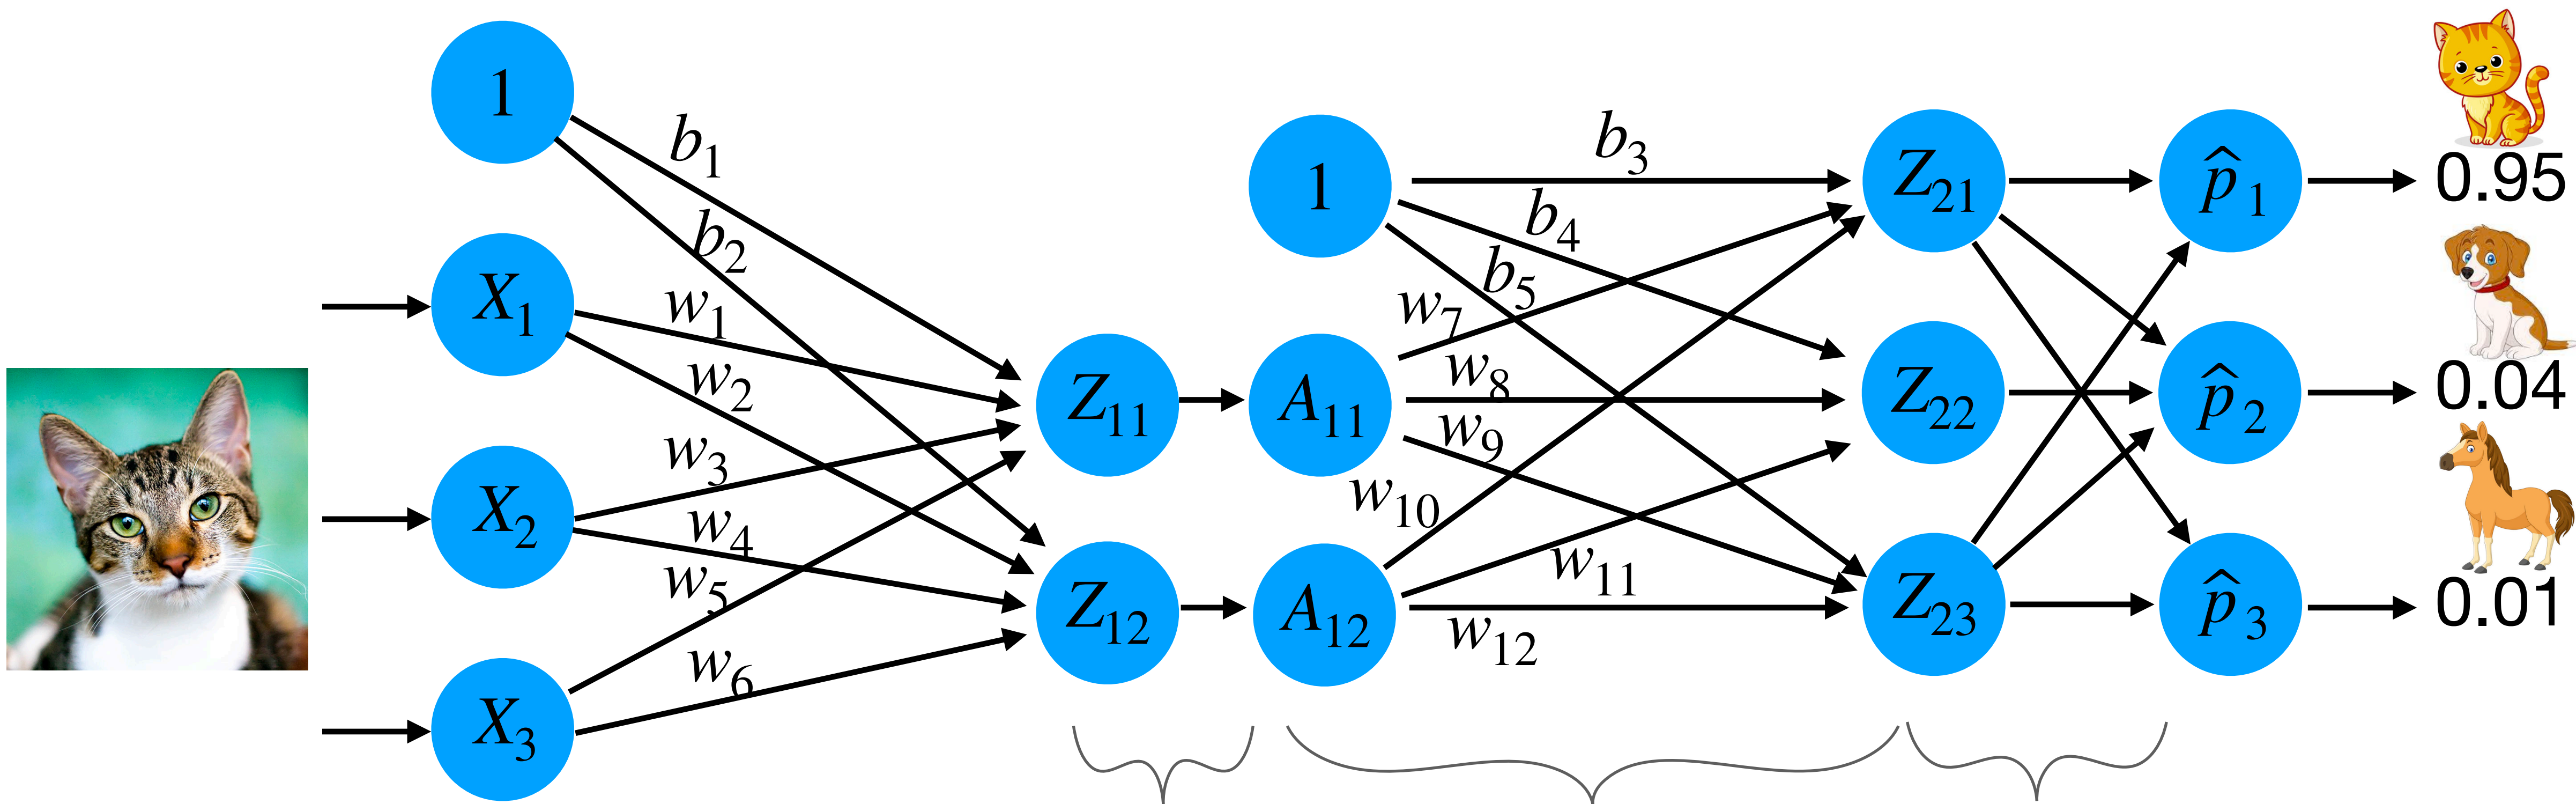


# One-hidden-layer neural network

Input layer

Hidden layer

Output layer



summation

ReLU

activation


summation

softmax

activation

E.g.,  $Z_{12} = b_2 + w_2X_1 + w_4X_2 + w_6X_3$

E.g.,  $A_{12} = \max(0, Z_{12})$

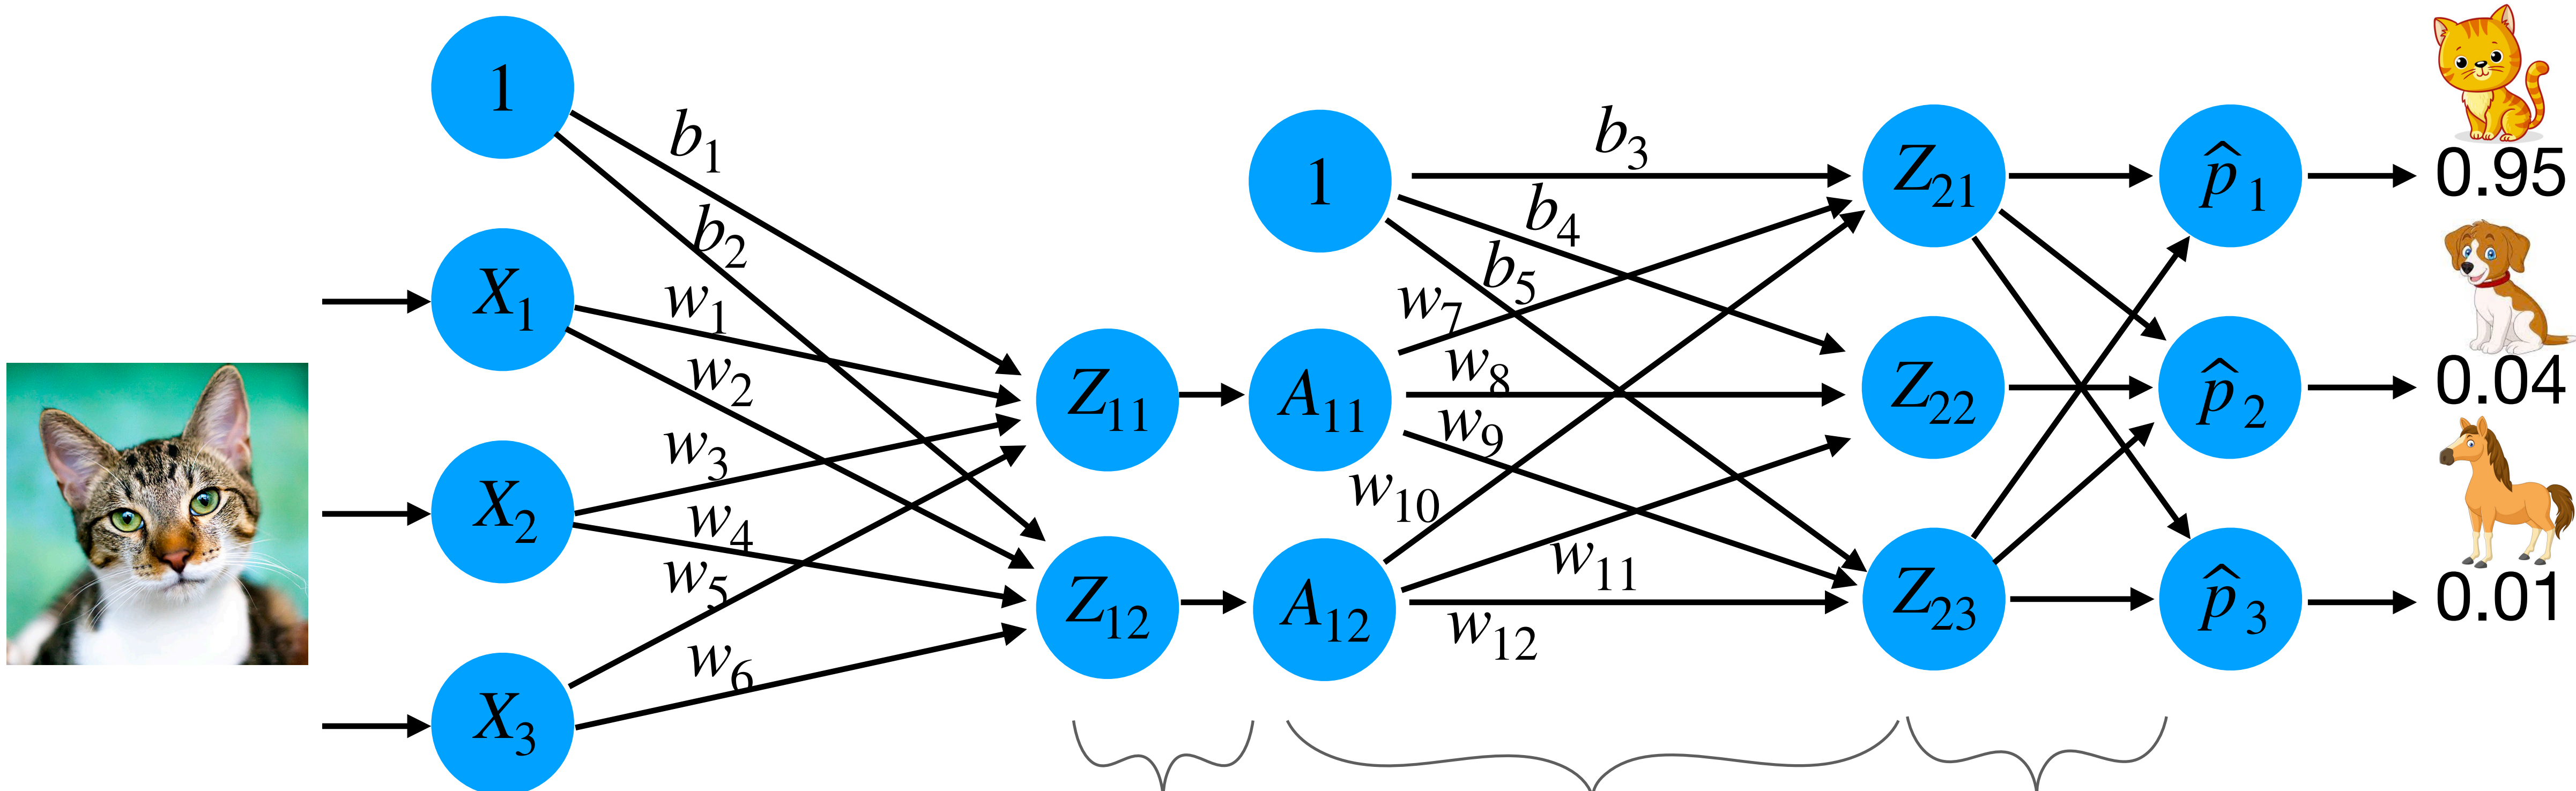


# One-hidden-layer neural network

Input layer

Hidden layer

Output layer



summation

E.g.,  $Z_{12} = b_2 + w_2X_1 + w_4X_2 + w_6X_3$

ReLU


activation

E.g.,  $A_{12} = \max(0, Z_{12})$

summation

softmax

activation



Terminology:


$b_j$ : biases

$w_j$ : weights



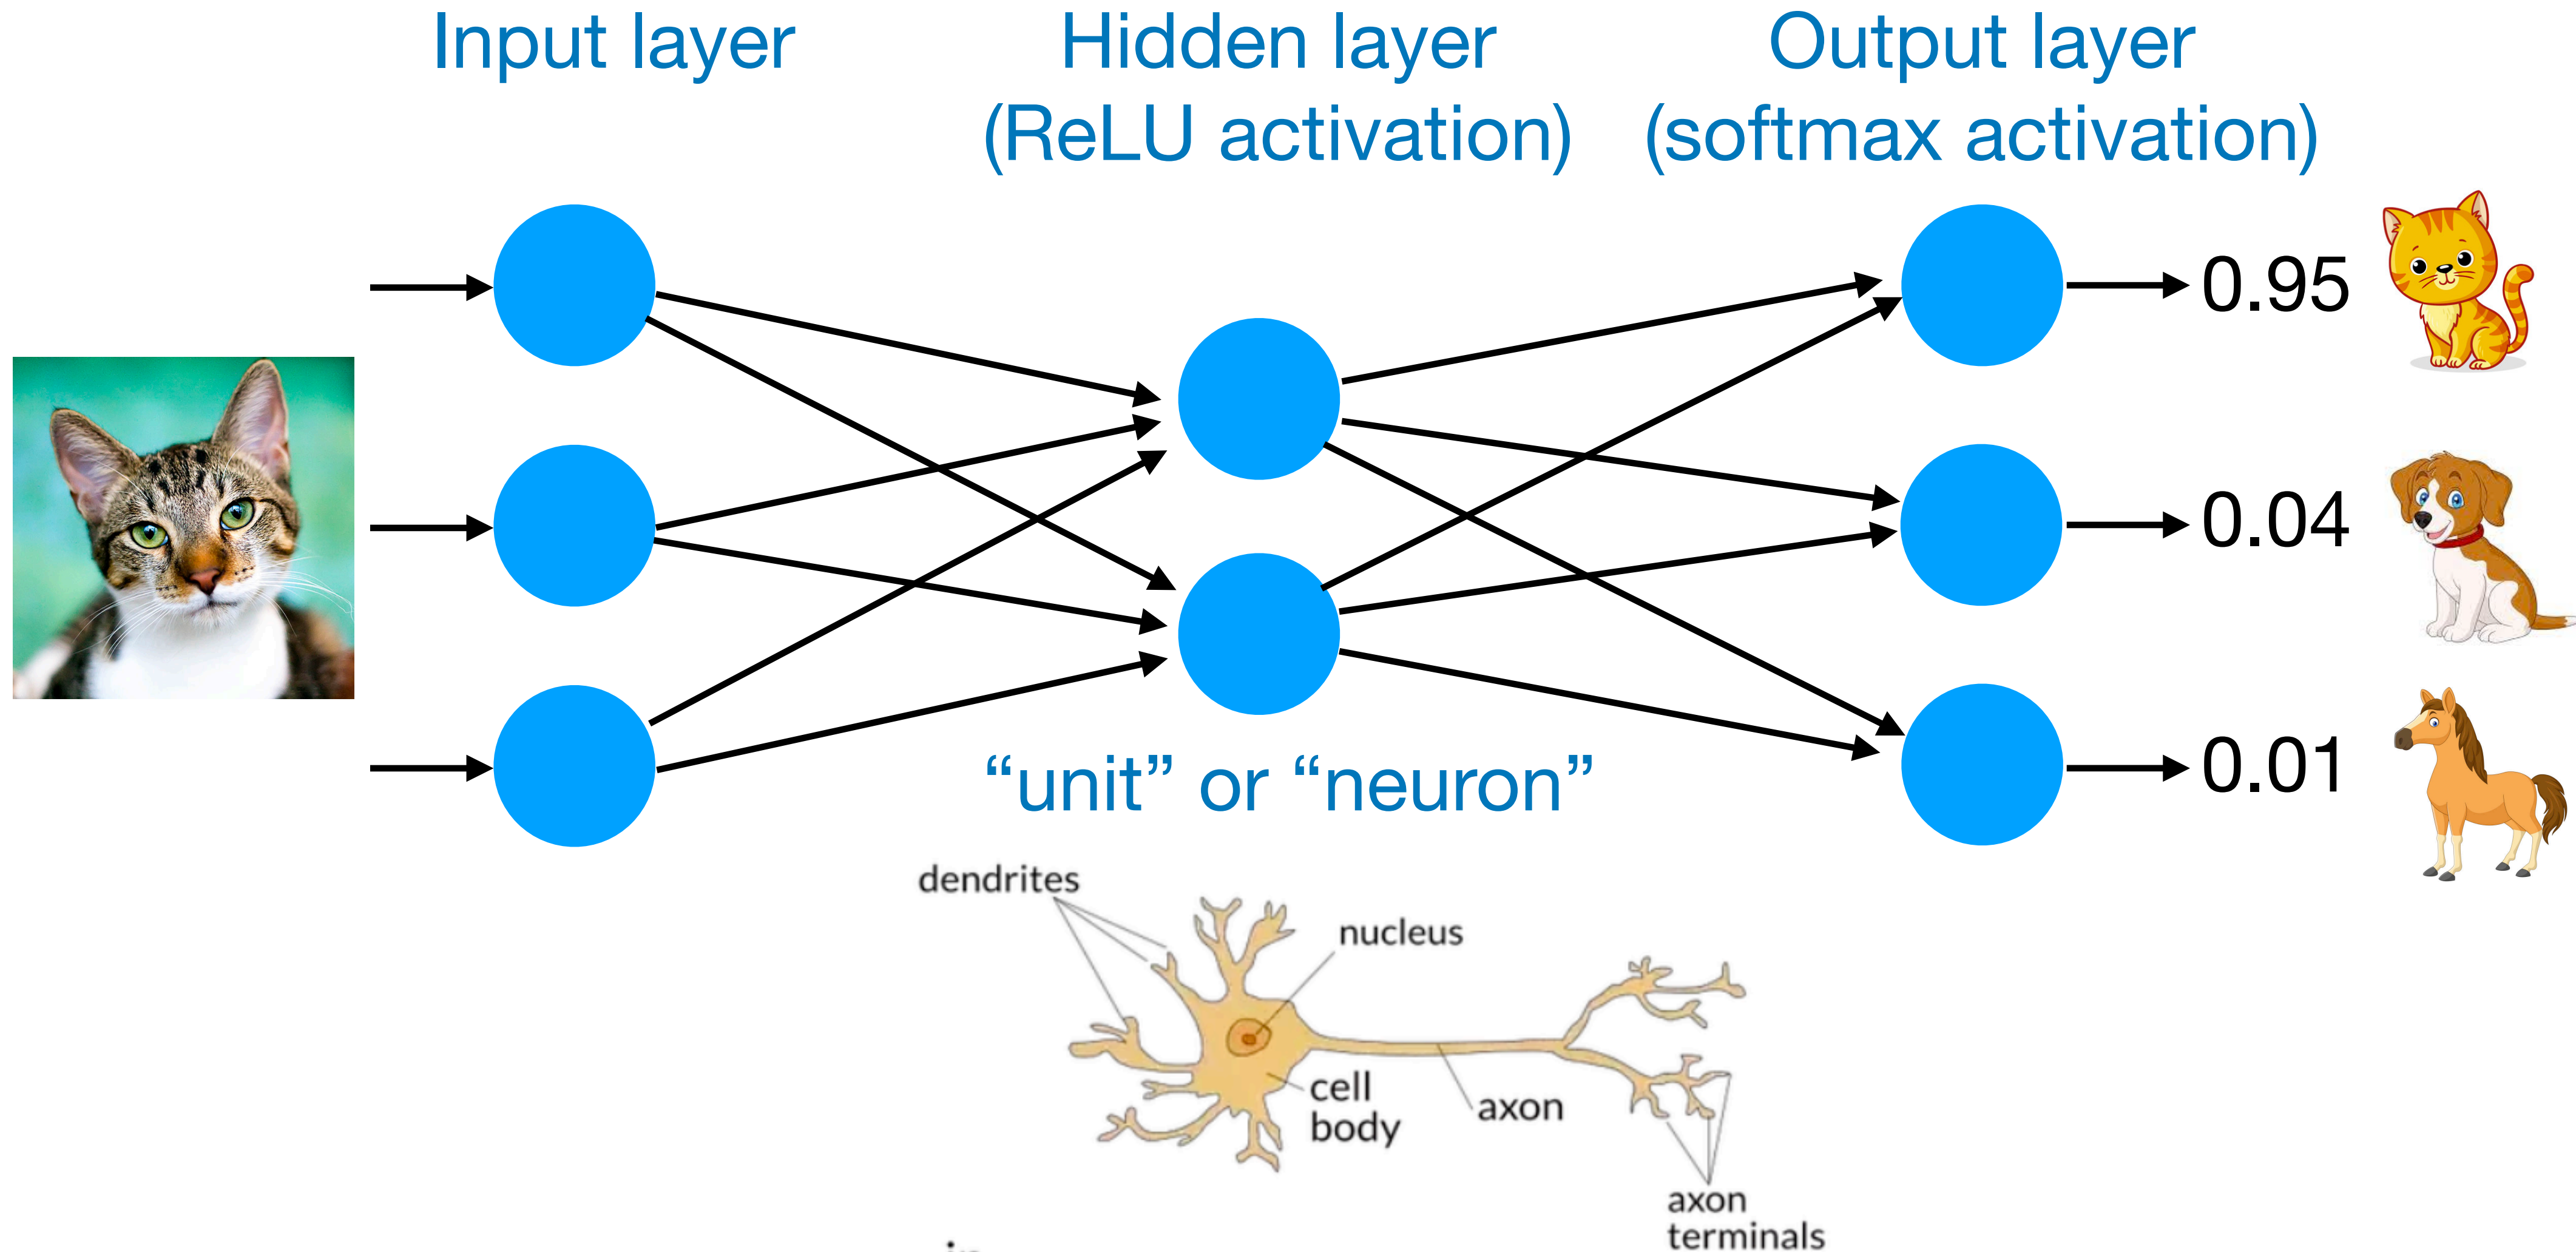
# One-hidden-layer neural network

## Concise representation



# One-hidden-layer neural network

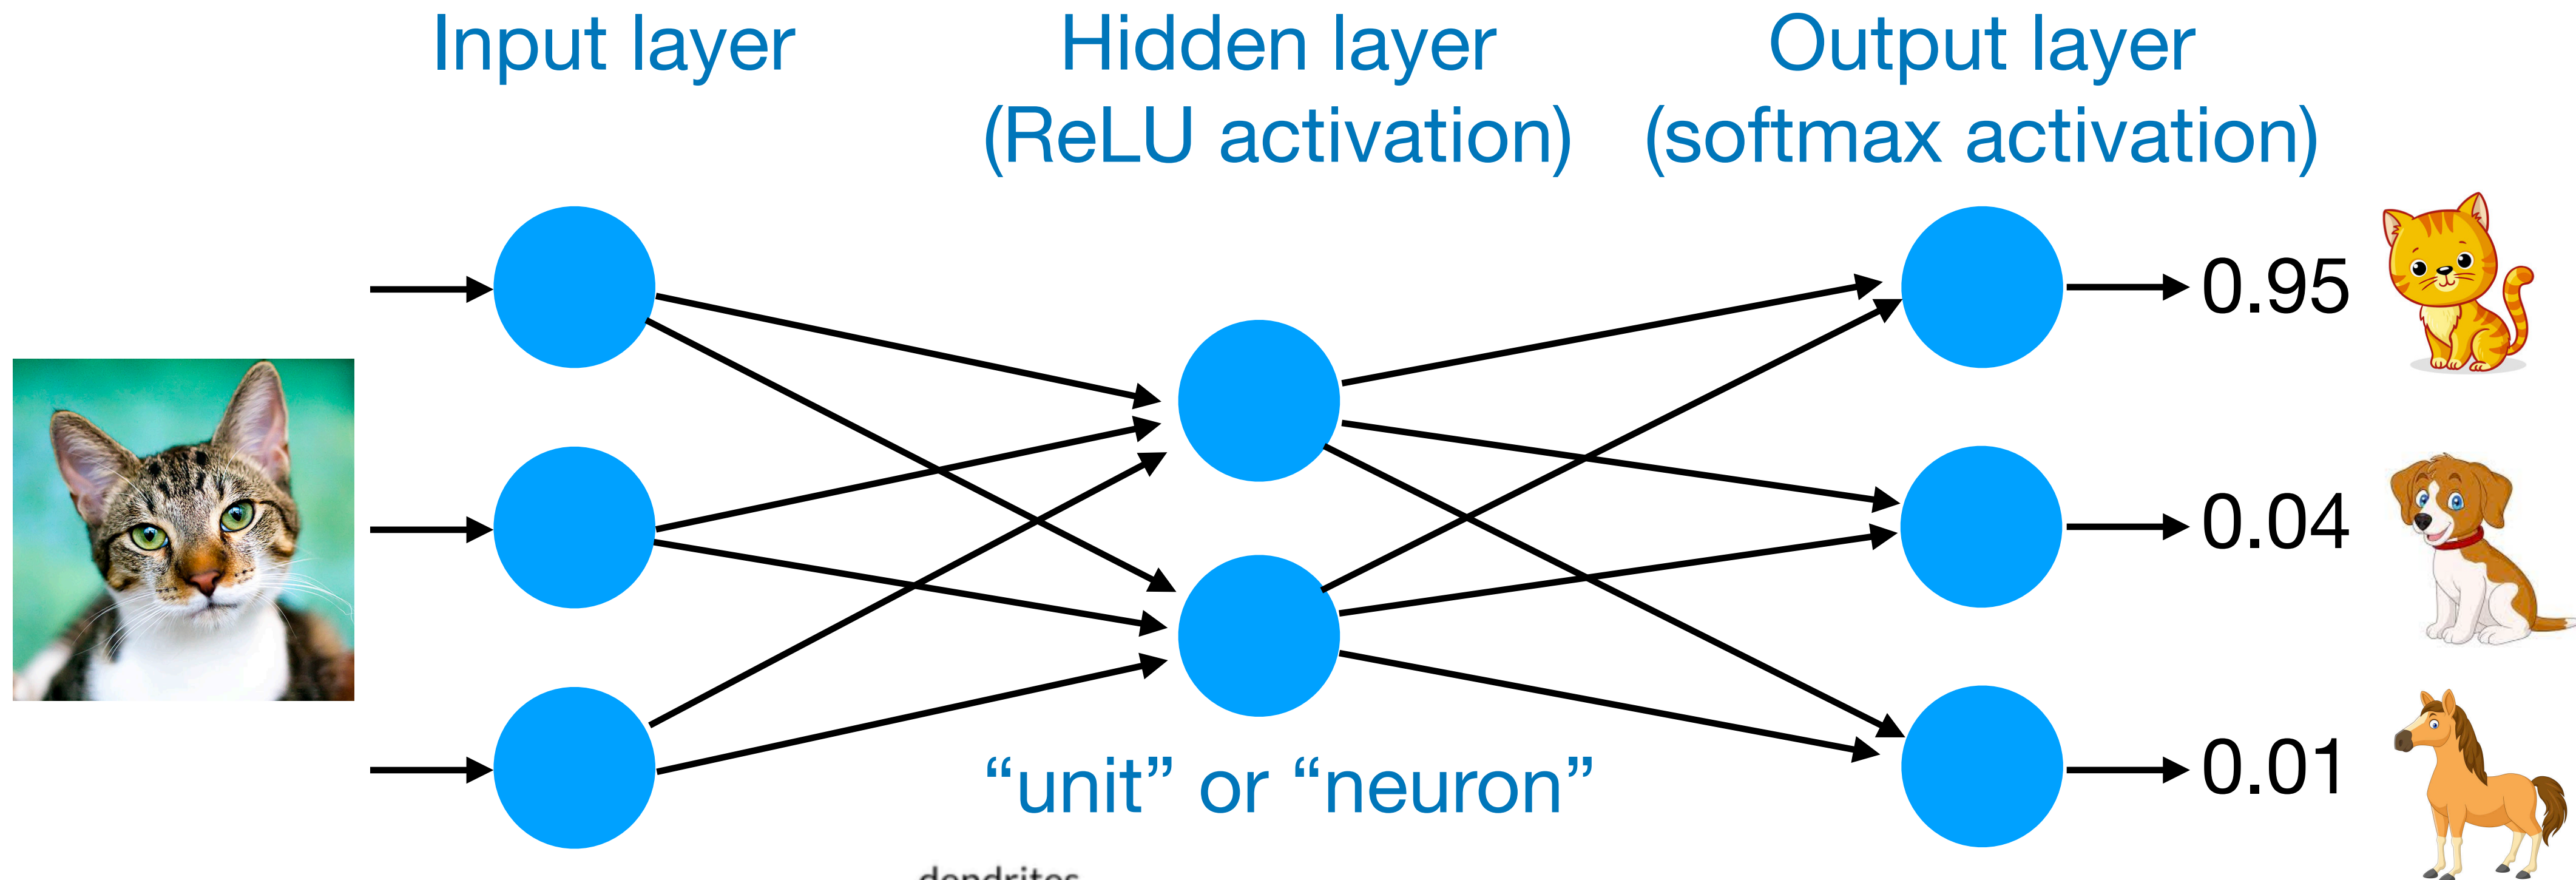
## Concise representation






# One-hidden-layer neural network

## Concise representation



### Notes:


- Bias terms left implicit.
- Fully connected architecture.
- Hidden layer: learned features.






# One-hidden-layer neural network

Even more concise representation



# One-hidden-layer neural network


Even more concise representation



Number of parameters in FC layer =  
 $(\# \text{ units in previous layer} + 1) \times (\# \text{ units in this layer})$

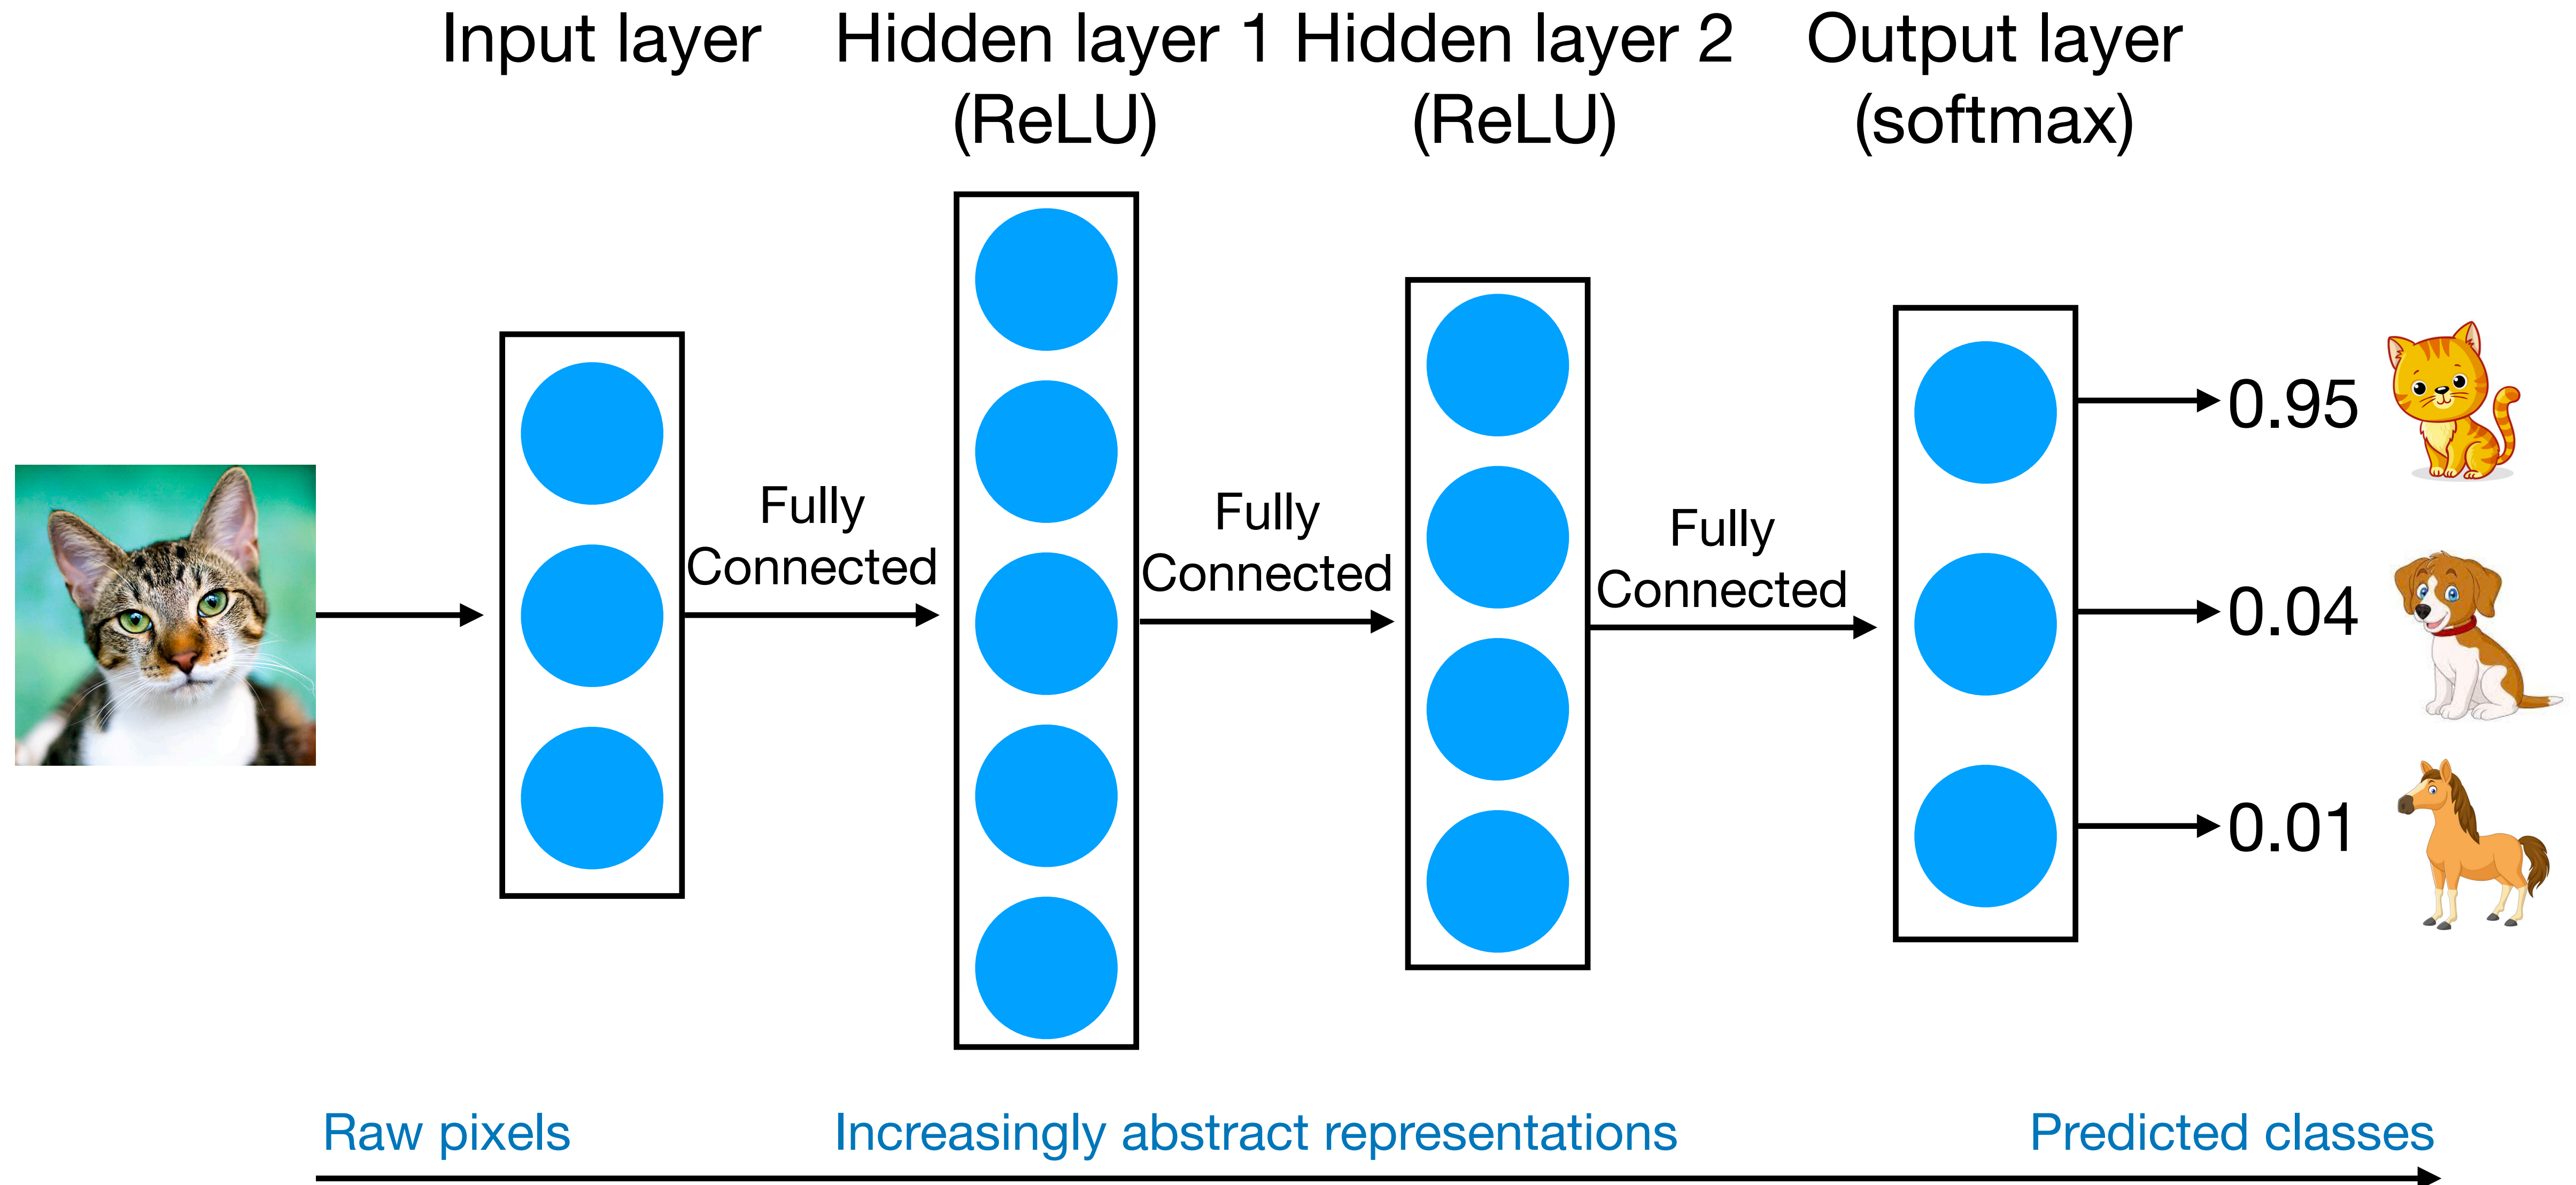
# Multi-layer fully connected neural networks

Input layer    Hidden layer 1 (ReLU)    Hidden layer 2 (ReLU)    Output layer (softmax)





# Multi-layer fully connected neural networks



# Training deep learning models

# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:

$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta) .$$




# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:

$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.



<https://arxiv.org/abs/1712.09913>


# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>






# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>






# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>




# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>






# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>




# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>






# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>




# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>






# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>






# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>




# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>






# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>






# Training deep learning models

Training a neural network amounts to solving our usual optimization problem:


$$\hat{\beta} = \arg \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)) + \lambda \cdot \text{penalty}(\beta) = \arg \min_{\beta} F(\beta).$$

Unfortunately, **neural networks have non-convex objective functions**, causing issues discussed in the previous lecture.

In practice, deep learning models are trained using **stochastic gradient descent**, which is a variant of gradient descent.



<https://arxiv.org/abs/1712.09913>



# The computational cost of gradient descent

# The computational cost of gradient descent

computational cost = (cost of computing the gradient)  $\times$  (number of iterations)



# The computational cost of gradient descent

computational cost = (cost of computing the gradient)  $\times$  (number of iterations)

Number of iterations has to do with the size of the learning rate (like boosting) and the shape of the objective function  $F$ .

# The computational cost of gradient descent

computational cost = (cost of computing the gradient)  $\times$  (number of iterations)

Number of iterations has to do with the size of the learning rate (like boosting) and the shape of the objective function  $F$ .

Cost of computing the gradient has to do with the size of the training data:

# The computational cost of gradient descent

computational cost = (cost of computing the gradient)  $\times$  (number of iterations)

Number of iterations has to do with the size of the learning rate (like boosting) and the shape of the objective function  $F$ .

Cost of computing the gradient has to do with the size of the training data:

$$\text{If } F(\beta) = \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)), \text{ then } \nabla F(\beta) = \frac{1}{n} \sum_{i=1}^n \nabla L(Y_i, f_{\beta}(X_i)).$$



# The computational cost of gradient descent

computational cost = (cost of computing the gradient)  $\times$  (number of iterations)

Number of iterations has to do with the size of the learning rate (like boosting) and the shape of the objective function  $F$ .

Cost of computing the gradient has to do with the size of the training data:

$$\text{If } F(\beta) = \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)), \text{ then } \nabla F(\beta) = \frac{1}{n} \sum_{i=1}^n \nabla L(Y_i, f_{\beta}(X_i)).$$

Each evaluation of the gradient requires a pass through the entire training data, and modern data sets can have millions of training observations.

# The computational cost of gradient descent

computational cost = (cost of computing the gradient)  $\times$  (number of iterations)

Number of iterations has to do with the size of the learning rate (like boosting) and the shape of the objective function  $F$ .

Cost of computing the gradient has to do with the size of the training data:

$$\text{If } F(\beta) = \frac{1}{n} \sum_{i=1}^n L(Y_i, f_{\beta}(X_i)), \text{ then } \nabla F(\beta) = \frac{1}{n} \sum_{i=1}^n \nabla L(Y_i, f_{\beta}(X_i)).$$

Each evaluation of the gradient requires a pass through the entire training data, and modern data sets can have millions of training observations.

This can make gradient descent prohibitively expensive.

# Stochastic gradient descent



# Stochastic gradient descent

Use subset  $S \subseteq \{1, \dots, n\}$  (mini-batch) of observations to approximate gradient:

$$\nabla F(\beta) = \frac{1}{n} \sum_{i=1}^n \nabla L(Y_i, f_{\beta}(X_i)) \approx \frac{1}{|S|} \sum_{i \in S} \nabla L(Y_i, f_{\beta}(X_i)).$$

# Stochastic gradient descent

Use subset  $S \subseteq \{1, \dots, n\}$  (mini-batch) of observations to approximate gradient:

$$\nabla F(\beta) = \frac{1}{n} \sum_{i=1}^n \nabla L(Y_i, f_{\beta}(X_i)) \approx \frac{1}{|S|} \sum_{i \in S} \nabla L(Y_i, f_{\beta}(X_i)).$$

1. Choose some initial value of  $\beta$ .

# Stochastic gradient descent

Use subset  $S \subseteq \{1, \dots, n\}$  (mini-batch) of observations to approximate gradient:

$$\nabla F(\beta) = \frac{1}{n} \sum_{i=1}^n \nabla L(Y_i, f_{\beta}(X_i)) \approx \frac{1}{|S|} \sum_{i \in S} \nabla L(Y_i, f_{\beta}(X_i)).$$

1. Choose some initial value of  $\beta$ .
2. Randomly assign observations to mini-batches  $S_1, \dots, S_M$  of a certain size



# Stochastic gradient descent

Use subset  $S \subseteq \{1, \dots, n\}$  (mini-batch) of observations to approximate gradient:

$$\nabla F(\beta) = \frac{1}{n} \sum_{i=1}^n \nabla L(Y_i, f_{\beta}(X_i)) \approx \frac{1}{|S|} \sum_{i \in S} \nabla L(Y_i, f_{\beta}(X_i)).$$

1. Choose some initial value of  $\beta$ .
2. Randomly assign observations to mini-batches  $S_1, \dots, S_M$  of a certain size
3. Repeat until convergence:

# Stochastic gradient descent

Use subset  $S \subseteq \{1, \dots, n\}$  (mini-batch) of observations to approximate gradient:

$$\nabla F(\beta) = \frac{1}{n} \sum_{i=1}^n \nabla L(Y_i, f_{\beta}(X_i)) \approx \frac{1}{|S|} \sum_{i \in S} \nabla L(Y_i, f_{\beta}(X_i)).$$

1. Choose some initial value of  $\beta$ .
2. Randomly assign observations to mini-batches  $S_1, \dots, S_M$  of a certain size
3. Repeat until convergence:

- For  $m = 1, \dots, M$ , update  $\beta \leftarrow \beta - \gamma \cdot \frac{1}{|S_m|} \sum_{i \in S_m} \nabla L(Y_i, f_{\beta}(X_i)).$

# Stochastic gradient descent

Use subset  $S \subseteq \{1, \dots, n\}$  (mini-batch) of observations to approximate gradient:

$$\nabla F(\beta) = \frac{1}{n} \sum_{i=1}^n \nabla L(Y_i, f_{\beta}(X_i)) \approx \frac{1}{|S|} \sum_{i \in S} \nabla L(Y_i, f_{\beta}(X_i)).$$

1. Choose some initial value of  $\beta$ .
2. Randomly assign observations to mini-batches  $S_1, \dots, S_M$  of a certain size
3. Repeat until convergence:

- $\underbrace{\text{For } m = 1, \dots, M, \text{ update } \beta \leftarrow \beta - \gamma \cdot \frac{1}{|S_m|} \sum_{i \in S_m} \nabla L(Y_i, f_{\beta}(X_i))}_{\text{One "epoch"}}$



# Stochastic gradient descent

Use subset  $S \subseteq \{1, \dots, n\}$  (mini-batch) of observations to approximate gradient:

$$\nabla F(\beta) = \frac{1}{n} \sum_{i=1}^n \nabla L(Y_i, f_{\beta}(X_i)) \approx \frac{1}{|S|} \sum_{i \in S} \nabla L(Y_i, f_{\beta}(X_i)).$$

1. Choose some initial value of  $\beta$ .
2. Randomly assign observations to mini-batches  $S_1, \dots, S_M$  of a certain size
3. Repeat until convergence:

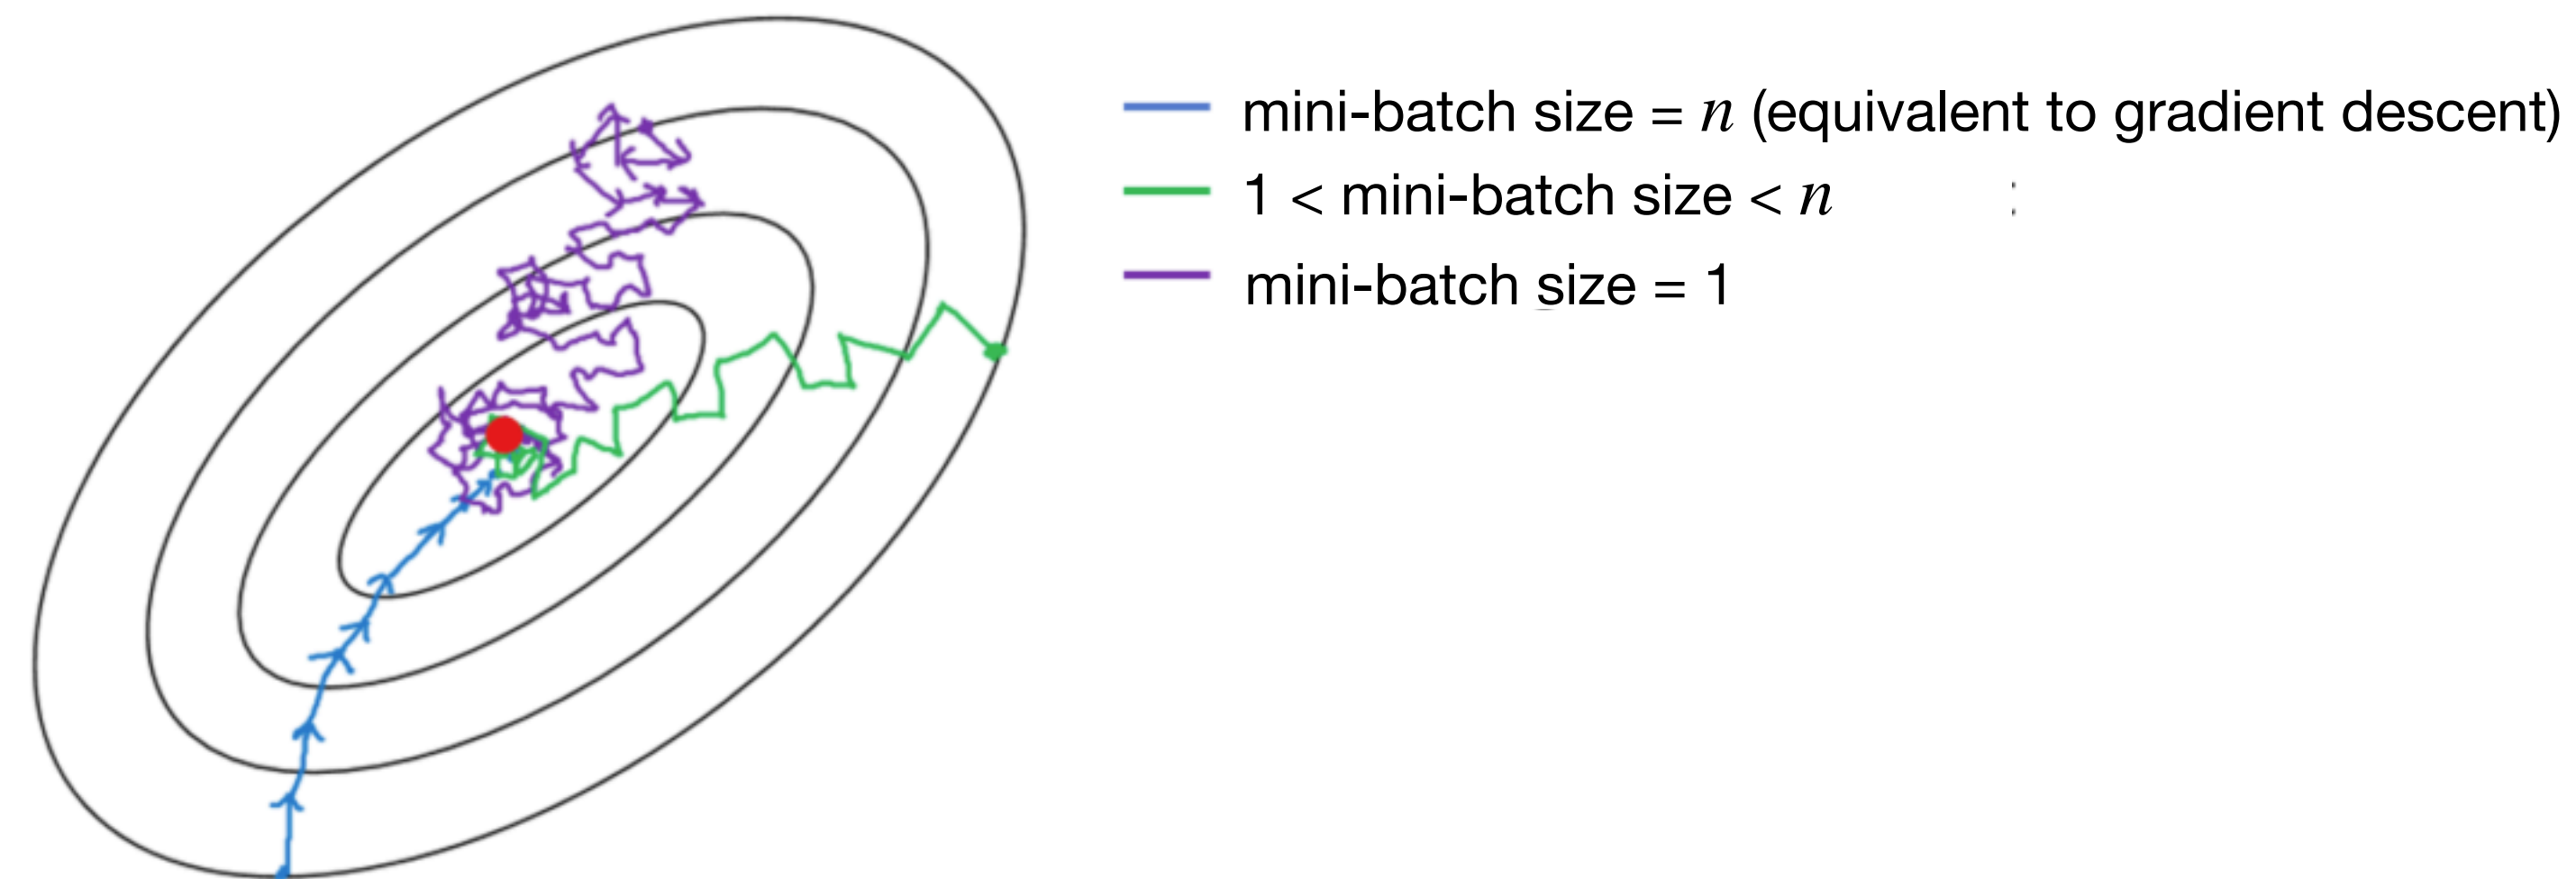
- For  $m = 1, \dots, M$ , update  $\beta \leftarrow \beta - \gamma \cdot \frac{1}{|S_m|} \sum_{i \in S_m} \nabla L(Y_i, f_{\beta}(X_i))$ .

One “epoch”

**Backpropagation:** An efficient algorithm to compute  $\nabla L(Y_i, f_{\beta}(X_i))$

# Behavior of stochastic gradient descent

Stochastic gradient descent wobbles toward decreasing values of the objective:




Source: <https://medium.com/analytics-vidhya/gradient-descent-vs-stochastic-gd-vs-mini-batch-sgd-fbd3a2cb4ba4>

The smaller the mini-batch, the cheaper and more wobbly each step is; Intermediate mini-batch sizes tend to work well, e.g. mini-batch size = 32.

**Bonus:** The extra randomness sometimes allows SGD to wobble past local minima.

# Stochastic gradient descent step versus epoch

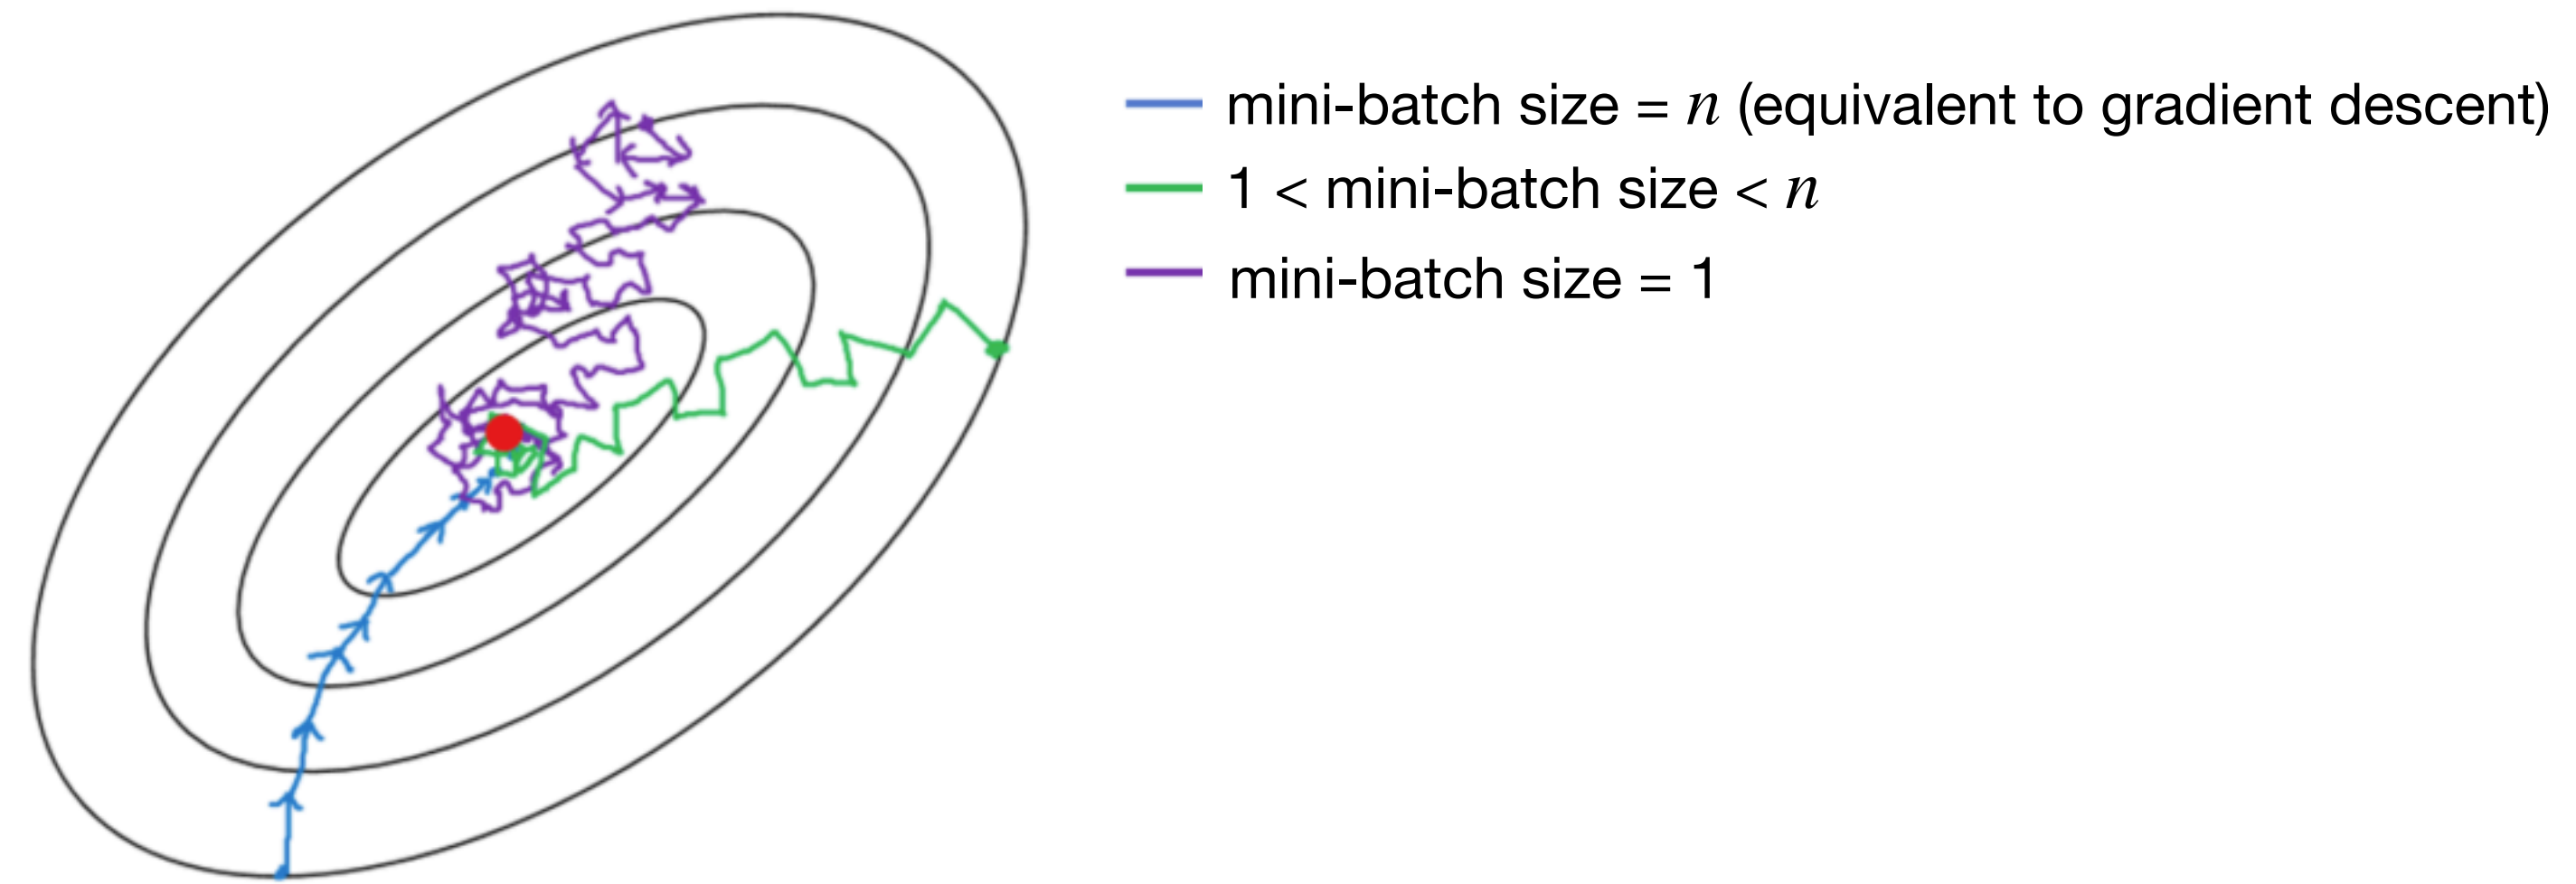


Source: <https://medium.com/analytics-vidhya/gradient-descent-vs-stochastic-gd-vs-mini-batch-sgd-fbd3a2cb4ba4>



# Stochastic gradient descent step versus epoch


An epoch consists of SGD steps that cycle through all of the observations once.



Source: <https://medium.com/analytics-vidhya/gradient-descent-vs-stochastic-gd-vs-mini-batch-sgd-fbd3a2cb4ba4>

# Stochastic gradient descent step versus epoch


An epoch consists of SGD steps that cycle through all of the observations once.



Source: <https://medium.com/analytics-vidhya/gradient-descent-vs-stochastic-gd-vs-mini-batch-sgd-fbd3a2cb4ba4>

# Stochastic gradient descent step versus epoch

An epoch consists of SGD steps that cycle through all of the observations once.




Source: <https://medium.com/analytics-vidhya/gradient-descent-vs-stochastic-gd-vs-mini-batch-sgd-fbd3a2cb4ba4>



# Stochastic gradient descent step versus epoch

An epoch consists of SGD steps that cycle through all of the observations once.




Source: <https://medium.com/analytics-vidhya/gradient-descent-vs-stochastic-gd-vs-mini-batch-sgd-fbd3a2cb4ba4>

Observation	1	2	3	4	5	6	7	8	9	1	2	3	4	5	6	7	8	9
Epoch	Epoch 1									Epoch 2								

# Stochastic gradient descent step versus epoch

An epoch consists of SGD steps that cycle through all of the observations once.




Source: <https://medium.com/analytics-vidhya/gradient-descent-vs-stochastic-gd-vs-mini-batch-sgd-fbd3a2cb4ba4>

Observation	1	2	3	4	5	6	7	8	9	1	2	3	4	5	6	7	8	9
Epoch	Epoch 1									Epoch 2								
SGD step (mini-batch size 1)	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18

# Stochastic gradient descent step versus epoch

An epoch consists of SGD steps that cycle through all of the observations once.




Source: <https://medium.com/analytics-vidhya/gradient-descent-vs-stochastic-gd-vs-mini-batch-sgd-fbd3a2cb4ba4>

Observation	1	2	3	4	5	6	7	8	9	1	2	3	4	5	6	7	8	9
Epoch	Epoch 1									Epoch 2								
SGD step (mini-batch size 1)	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
SGD step (mini-batch size 3)	Step 1			Step 2			Step 3			Step 4			Step 5			Step 6		



# Stochastic gradient descent step versus epoch

An epoch consists of SGD steps that cycle through all of the observations once.




Source: <https://medium.com/analytics-vidhya/gradient-descent-vs-stochastic-gd-vs-mini-batch-sgd-fbd3a2cb4ba4>

E.g.  $n = 9$ , mini-batch size 3. Then each epoch consists of three SGD steps.

Observation	1	2	3	4	5	6	7	8	9	1	2	3	4	5	6	7	8	9
Epoch	Epoch 1									Epoch 2								
SGD step (mini-batch size 1)	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
SGD step (mini-batch size 3)	Step 1			Step 2			Step 3			Step 4			Step 5			Step 6		

# Stochastic gradient descent step versus epoch

An epoch consists of SGD steps that cycle through all of the observations once.




Source: <https://medium.com/analytics-vidhya/gradient-descent-vs-stochastic-gd-vs-mini-batch-sgd-fbd3a2cb4ba4>

E.g.  $n = 9$ , mini-batch size 3. Then each epoch consists of three SGD steps.

Observation	1	2	3	4	5	6	7	8	9	1	2	3	4	5	6	7	8	9
Epoch	Epoch 1									Epoch 2								
SGD step (mini-batch size 1)	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
SGD step (mini-batch size 3)	Step 1			Step 2			Step 3			Step 4			Step 5			Step 6		
Gradient step	Step 1									Step 2								


# The learning rate for (stochastic) gradient descent



Source: <https://www.jeremyjordan.me/nn-learning-rate/>




# The learning rate for (stochastic) gradient descent



Source: <https://www.jeremyjordan.me/nn-learning-rate/>

- Setting the learning rate is more of an art than a science; might need to try a few values to get a good one.

# The learning rate for (stochastic) gradient descent



Source: <https://www.jeremyjordan.me/nn-learning-rate/>

- Setting the learning rate is more of an art than a science; might need to try a few values to get a good one.
- Especially for non-convex optimization, people come up with clever strategies like shrinking learning rates, cycling learning rates, adaptive learning rates, etc. (RMSprop, Adam, AdaGrad, AdaDelta, ...)

# Regularization in deep learning models



# Regularization in deep learning models

Explicit regularization via penalization. Ridge regression penalty is the most common; this kind of penalization is known as weight decay.

# Regularization in deep learning models

**Explicit regularization** via penalization. Ridge regression penalty is the most common; this kind of penalization is known as **weight decay**.


**Implicit regularization:** Other techniques to control complexity of the model.

# Regularization in deep learning models

**Explicit regularization** via penalization. Ridge regression penalty is the most common; this kind of penalization is known as **weight decay**.

**Implicit regularization:** Other techniques to control complexity of the model.

- **Early stopping:** Do not run SGD until convergence; rather, stop SGD when validation error starts to increase.







# Regularization in deep learning models

**Explicit regularization** via penalization. Ridge regression penalty is the most common; this kind of penalization is known as **weight decay**.

**Implicit regularization:** Other techniques to control complexity of the model.

- **Early stopping:** Do not run SGD until convergence; rather, stop SGD when validation error starts to increase.
- **Dropout:** At each SGD iteration, remove a randomly selected set of nodes from the network (analogous to sub-sampling features for random forests).



# Model complexity and tuning in deep learning

# Model complexity and tuning in deep learning

Deep learning has many tuning parameters:



# Model complexity and tuning in deep learning

Deep learning has many tuning parameters:

- Network architecture (number of layers, units per layer)

# Model complexity and tuning in deep learning

Deep learning has many tuning parameters:

- Network architecture (number of layers, units per layer)
- Optimization algorithm used (RMSprop, Adam, etc.)

# Model complexity and tuning in deep learning

Deep learning has many tuning parameters:

- Network architecture (number of layers, units per layer)
- Optimization algorithm used (RMSprop, Adam, etc.)
- Weight decay parameter, dropout rate, number of SGD iterations



# Model complexity and tuning in deep learning

Deep learning has many tuning parameters:

- Network architecture (number of layers, units per layer)
- Optimization algorithm used (RMSprop, Adam, etc.)
- Weight decay parameter, dropout rate, number of SGD iterations

More complex models arise from bigger networks, less weight decay, less dropout, more SGD iterations.

# Model complexity and tuning in deep learning

Deep learning has many tuning parameters:

- Network architecture (number of layers, units per layer)
- Optimization algorithm used (RMSprop, Adam, etc.)
- Weight decay parameter, dropout rate, number of SGD iterations

More complex models arise from bigger networks, less weight decay, less dropout, more SGD iterations.

Deep learning performance can be sensitive to these parameters; there is no standard way to tune neural networks; tuning is computationally expensive.

# Model complexity and tuning in deep learning

Deep learning has many tuning parameters:

- Network architecture (number of layers, units per layer)
- Optimization algorithm used (RMSprop, Adam, etc.)
- Weight decay parameter, dropout rate, number of SGD iterations

More complex models arise from bigger networks, less weight decay, less dropout, more SGD iterations.

Deep learning performance can be sensitive to these parameters; there is no standard way to tune neural networks; tuning is computationally expensive.

For computational savings, validation sets are typically employed instead of cross-validation to tune neural networks.

# Summary

- Adding hidden layers to logistic regression gives multilayer neural networks.
- Multilayer neural networks build increasingly complex representations of the input data (feature learning).
- Stochastic gradient descent is used to train neural networks; it uses noisy approximations to the gradient but is much faster than gradient descent.
- Explicit and implicit regularization are used to control model complexity, the latter including early stopping and dropout.
- Neural networks come with a smorgasbord of tuning parameters, which often require considerable effort to tune.



# Summary

- Adding hidden layers to logistic regression gives multilayer neural networks.
- Multilayer neural networks build increasingly complex representations of the input data ([feature learning](#)).
- [Stochastic gradient descent](#) is used to train neural networks; it uses noisy approximations to the gradient but is much faster than gradient descent.
- Explicit and implicit regularization are used to control model complexity, the latter including [early stopping](#) and [dropout](#).
- Neural networks come with a smorgasbord of tuning parameters, which often require considerable effort to tune.

[Quiz Practice](#)